

Introduction to algorithm

Eric's Notes

Eric

February 10, 2014

Contents

1	课程简介及算法分析	5
1.1	Insertion sort	5
1.2	Merge sort	6
1.3	Recursion tree	8
2	渐进符号, 递归及解法	8
2.1	Standard notations and common functions	10
2.2	Substitution method	11
2.2.1	Subtleties	11
2.2.2	Changing variables	11
2.3	Recursion-tree method	11
2.4	Master method	12
2.4.1	Proof of master method	13
3	分治法 divide and conquer	14
3.1	Merge sort	14
3.2	Binary search	14
3.3	Powering a number	15
3.4	Fibonacci numbers	15
3.5	Matrix multiplication	15
4	快排及随机算法	17
4.1	Analysis	18
4.2	Random quicksort	20
4.3	Median-of-3 Pivot	21
4.4	Random method analysis	21
4.5	Exercise	23
5	Heapsort	23
5.1	Maintaining the heap property	24
5.2	Building a heap	25
5.3	The heapsort algorithm	26
5.4	Priority queues	26
6	线性时间排序	27
6.1	decision tree model 决策树	27
6.2	Counting sort 计数排序	28
6.3	Radix sort 基数排序	29
6.3.1	说明	29
6.3.2	Analysis	30
7	Order statistics	31

8 Hashing	31
8.1 Direct address tables	32
8.2 Hash tables	32
8.3 Hash functions	34
8.3.1 Division method	34
8.3.2 Multiplication method	34
8.3.3 Collision resolution by open addressing	35
9 Universal hashing and perfect hashing	38
9.1 Universal hashing	38
9.2 Perfect hashing 完全哈希	39
9.2.1 Collision-free	39
9.2.2 Analysis of storage	40
10 Binary search tree(BST)	41
10.1 what is a binary search tree?	41
10.1.1 tree walk 遍历	41
10.2 Querying a binary search tree	41
10.2.1 Searching	42
10.2.2 Minimum and maximum	42
10.2.3 Successor and predecessor	42
10.3 Insertion and deletion	43
10.4 Randomly built binary search trees	46
11 Red-Black Tree	46
11.1 Properties of red-black trees	46
11.2 Rotations	47
11.3 Insertion	48
11.4 Deletion	49
11.5 Analysis	51
12 Augmenting data structures	52
12.1 Dynamic order statistics	52
12.2 How to augment a data structure	53
12.3 Interval tree	53
13 Skip list	55
14 Dynamic programming	55
14.1 Assembly-line scheduling	56
14.2 Longest common subsequence (LCS)	58
14.3 Fibonacci	59
14.4 Optimal binary search trees	59
14.5 Common subproblems	59

15 贪婪算法和最小生成树	60
15.1 Greedy versus dynamic programming	60
15.2 Tree	60
15.3 Huffman code	61
15.4 MST - Minimum spanning tress	62
16 Amortized Analysis	62
16.1 Incrementing a binary counter	62
16.2 Dynamic tables	63
16.3 Accounting method	64
16.4 Potential method	64
17 Advanced data structures	65
18 B-Trees	66
18.1 Definition of B-Trees	66
19 AVL tree	66
20 Graph	67
21 最短路径算法	67
21.1 Dijkstra's algorithm	68
22 String matching	69
23 并行	69
24 Appendix	70
24.1 Summations	70
24.1.1 Approximation by integrals	70
24.2 Relations	71
24.2.1 Functions	72
24.3 Graphs	72
24.3.1 Trees	74
24.4 Probability	76
24.4.1 Baye's theorem	77
24.4.2 Discrete random variables	78
24.4.3 Indicator random variable	79

1 课程简介及算法分析

1.1 Insertion sort

Pseudocode

INSERTION-SORT(A)

```

1 for j ← 2 to length[A]  #1是第一个元素
2   do key ← A[j] //将将要插入的数据保存下来
3   //Insert A[j] into the sorted sequence A[1 _ j - 1].
4   i ← j - 1
5   while i > 0 and A[i] > key
6     do A[i + 1] ← A[i]
7     i ← i - 1
8   A[i + 1] ← key

```




Figure 1: Insertion sort

ex:

```

8 2 4 9 3 6
2 8 4 9 3 6 // 2 goes before 8
2 4 8 9 3 6 // 4 goes before 8
2 4 8 9 3 6 // 9 stays there
2 3 4 8 9 6 // 3 goes before 4
2 3 4 6 8 9 // 6 goes before 8

```

python code

```

def insertionSort(L):
    #in place algorithm, 升序
    for j in range(1, len(L)):
        #将将要排序的数据保存下来, 开始对第j 个元素进行排序
        key=L[j]

```




Figure 2: Insertion sort

```

i=j-1
#找到存放key的位置
while i >= 0 and L[i]>key:
    L[i+1]=L[i] ## 将大数往后移动
    i=i-1
##循环结束,说明L[i] <= key,所以key应该放在i+1处
L[i+1]=key
return L

```

running time

1 already sorted: 最理想情况

2 reverse sorted: 最差情况

we want upper bounds 上界

kinds of analysis

worst-case $T(n) = \max$ time of any input of size n

average-case $T(n) = \text{expected time}$ 期望时间

(need assumption of statistical distribution of inputs)

best-case (bogus 假象)

BIG IDEA: **asymptotic analysis** 渐进分析

not the the exact running time of an algorithm

the order of growth of the running time

the insertion sort analysis

worst-case: input reverse sorted

$$T(n) = \sum_{j=2}^{j=n} \theta(j) = \theta(n^2) \text{ 算术级数 arithmetic series}$$

1.2 Merge sort

算法

$T(n)$ merge sort $A[1\dots n]$

$\theta(1)$ 1 if $n=1$, done

$2 \times \theta(n/2)$ 2 Recursively sort $A[1\dots \text{upper}(n/2)]$ and $A[\text{upper}(n/2)+1\dots n]$ 向上取整

$\theta(n)$ 3 merge 2 sorted list

Merge

Where is the smallest element of any two lists that are already sorted?

It is in one of two places, the head of the first list or the head of the second list

Key subroutine Merge

2 7 13 20

1 9 11 12

在两个 list head 中,1 最小, 所以 1 是 n 个元素中最小的, 排在最终的 list 的第一个位置, 现在总 list 和两个子 list 成为:

1

2 7 13 20

9 11 12

然后在比较两个子 list 中 head 位置那个更小, 把它放在总 list 的第二个位置

1 2

7 13 20

9 11 12

一直这么继续下去, 这里的每一步都是固定数目的操作, 和每一步中的数组的尺寸无关, 每一步总, 我们只关注两个 head, 并挑出最小的, 再把数组指针推进一位, 所以我知道当前的标头在哪里. 所以, 对于总数为 n 的输入, 时间是 $\theta(n)$ 的所以把两个 list 遍历和排序的时间是 $\theta(n)$, 有时我们称之为线性时间

merge sort 的例子

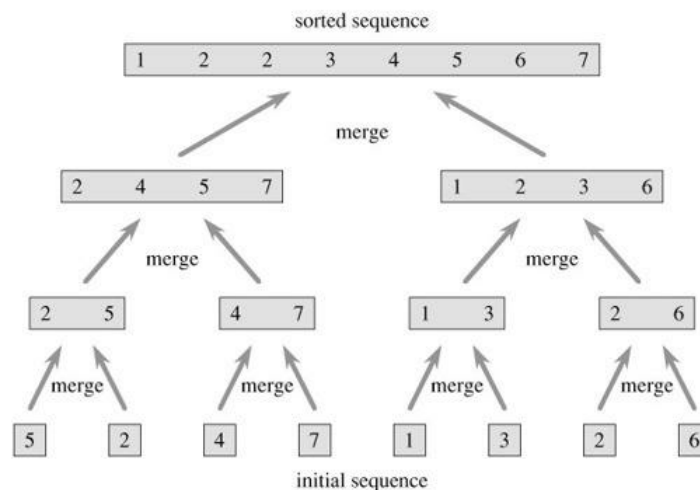


Figure 2.4: The operation of merge sort on the array $A = \langle 5, 2, 4, 7, 1, 3, 2, 6 \rangle$

Figure 3: Merge sort example

时间复杂度

$$T(n) = \begin{cases} \theta(1) & \text{si } n = 1 \\ 2 \times T(n/2) + \theta(n) & \text{si } n > 1 \end{cases}$$

1.3 Recursion tree

一直做下去, 得到, 每一行的和都为 cn , 数的深度为 $\lg n$, 树的 level 是 $\lg n + 1$, 最后一层的叶节点有 n 个, 每一个都是 $\theta(1)$




Figure 4: Recursion tree example

为了便于计算, 我们在这里设 $\theta(1) = c$; 先把每一层的加起来, 得到都是 cn , 然后再把所有层加起来, 得到

$$T(n) = (\lg n + 1)cn = \Theta(n \lg n)$$

So merge sort beats insertion sort.

$$\theta(n \lg n) < \theta(n^2)$$

2 渐进符号, 递归及解法

Définition 1. BIG O: 上界 \leq

$O(g(n)) = \{f(n) : \text{there exist positive constants } c \text{ and } n_0 \text{ such that } 0 \leq f(n) \leq cg(n) \text{ for all } n \geq n_0\}$

We say that $g(n)$ is an asymptotically upper bound for $f(n)$

BIG Ω 下界 \geq

$\Omega(g(n)) = \{f(n) : \text{there exist positive constants } c \text{ and } n_0 \text{ such that } 0 \leq cg(n) \leq f(n) \text{ for all } n \geq n_0\}$

We say that $g(n)$ is an asymptotically lower bound for $f(n)$.


Figure 3.1: Graphic examples of the Θ , O , and Ω notations. In each

Figure 5: Asymptotical notation

BIG Θ =

$\Theta(g(n)) = \{f(n) : \text{there exist positive constants } c_1, c_2, \text{ and } n_0 \text{ such that } 0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n) \text{ for all } n \geq n_0\}$

$$\Theta(g(n)) = O(g(n)) \cap \Omega(g(n))$$

We say that $g(n)$ is an asymptotically tight bound for $f(n)$.

The definition of $\Theta(g(n))$ requires that every member $f(n)$ in $\Theta(g(n))$ be asymptotically nonnegative, that is, that $f(n)$ be nonnegative whenever n is sufficiently large.

Ex: $2n^2 + \Theta(n) = \Theta(n^2)$ for any function $f(n)$ in $\Theta(n)$, there is some function $g(n)$ in $\Theta(n^2)$ such that $2n^2 + f(n) = g(n)$ for all n . In other words, the right-hand side of an equation provides a coarser level of detail than the left-hand side.

o -notation to denote an upper bound that is not asymptotically tight.

$o(g(n)) = \{f(n) : \text{for any positive constant } c > 0, \text{ there exists a constant } n_0 > 0 \text{ such that } 0 \leq f(n) < cg(n) \text{ for all } n \geq n_0\}$

For example, $2n = o(n^2)$, but $2n^2 \neq o(n^2)$

o -notation 表示的是一种相差比较大的

in the o -notation, the function $f(n)$ becomes insignificant relative to $g(n)$ as n approaches infinity; that is,

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0$$

ω -notation By analogy, ω -notation is to Ω -notation as o -notation is to O -notation.

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = \infty$$

Comparison of functions

与数的比较类比

- $f(n) = O(g(n)) \approx a \leq b$
- $f(n) = \Omega(g(n)) \approx a \geq b$
- $f(n) = \Theta(g(n)) \approx a = b$

- $f(n) = o(g(n)) \approx a \ll b$
- $f(n) = \omega(g(n)) \approx a \gg b$

2.1 Standard notations and common functions

函数的单调性

monotonically increasing 单调增

monotonically decreasing

strictly increasing 严格递增

strictly decreasing

Floors and ceilings

$$x - 1 < \lfloor x \rfloor \leq x \leq \lceil x \rceil < x + 1$$

For any integer n , $\lceil n/2 \rceil + \lfloor n/2 \rfloor = n$

And for any real number $n \geq 0$ and integer $a, b > 0$

1. $\lceil \lceil n/a \rceil / b \rceil = \lceil n/(ab) \rceil$
2. $\lfloor \lfloor n/a \rfloor / b \rfloor = \lfloor n/(ab) \rfloor$
3. $\lceil a/b \rceil \leq (a + (b - 1)) / b$
4. $\lfloor a/b \rfloor \geq (a - (b - 1)) / b$
5. $a \bmod n = a - \lfloor a/n \rfloor n$

$\lim_{n \rightarrow \infty} \frac{n^b}{a^n} = 0$ from which we can conclude that $n^b = o(a^n)$. Thus, any exponential function with a base strictly greater than 1 grows faster than any polynomial function.

$$\lim_{n \rightarrow \infty} \frac{\lg^b n}{(2^a)^{\lg n}} = \lim_{n \rightarrow \infty} \frac{\lg^b n}{n^a} = 0$$

$\lg^b n = o(n^a)$, for any constant $a > 0$. Thus, any positive polynomial function grows faster than any polylogarithmic function.

$$n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n (1 + \Theta(\frac{1}{n}))$$

Functional iteration

We use the notation $f^{(i)}(n)$ to denote the function $f(n)$ iteratively applied i times to an initial value of n . Formally, let $f(n)$ be a function over the reals. For nonnegative integers i , we recursively define

$$f^{(i)}(n) = \begin{cases} n & \text{si } i = 0 \\ f(f^{(i-1)}(n)) & \text{si } i > 0 \end{cases}$$

For example, if $f(n) = 2n$, then $f^{(i)}(n) = 2^i n$.

2.2 Substitution method

Substitution method for solving recurrences entails two steps:

Guess the form of the solution.

Use mathematical induction(数学归纳法) to find the constants and show that the solution works.

Ex: $T(n) = 2T(\lfloor n/2 \rfloor) + n$

猜测 $T(n) = O(n \lg n)$, 然后归纳法证明

$T(n) = 2T(n/2 + 17) + n$ 与上一式相差 17, 但是当 n 很大时, 17 可以忽略掉, 所以仍然猜测 $T(n) = O(n \lg n)$, 然后尝试用归纳法证明, 发现是正确的

2.2.1 Subtleties

$$T(n) = T(n/2) + T(n/2) + 1.$$

我们猜测 $T(n) \leq cn$

$T(n) \leq cn/2 + cn/2 + 1 = cn + 1$, wrong, 但是只差了一个常数

we're only off by the constant 1, a lower-order term, 加上一个 lower-order term, 猜测 $T(n) \leq cn - b$

$$T(n) \leq (cn/2 - b) + (cn/2 - b) + 1 = cn - 2b + 1 \leq cn - b$$

imply $b \geq 1$, 所以当 $b \geq 1$ 时, $T(n) \leq cn - b$

2.2.2 Changing variables

$$T(n) = 2T(\lfloor \sqrt{n} \rfloor) + \lg n$$

Renaming $m = \lg n$ yields $T(2^m) = 2T(2^{m/2}) + m$.

We can now rename $S(m) = T(2^m)$ to produce the new recurrence $S(m) = 2S(m/2) + m$,

这个见过, $S(m) = O(m \lg m)$.

Changing back from $S(m)$ to $T(n)$, we obtain

$$T(n) = T(2^m) = S(m) = O(m \lg m) = O(\lg n \lg \lg n).$$

2.3 Recursion-tree method

In a recursion tree, each node represents the cost of a single subproblem somewhere in the set of recursive function invocations.

We sum the costs within each level of the tree to obtain a set of per-level costs, and then we sum all the per-level costs to determine the total cost of all levels of the recursion.

Recursion trees are particularly useful when the recurrence describes the running time of a divide-and-conquer algorithm.

A recursion tree is best used to generate a good guess, which is then verified by the substitution method. So we can often tolerate a small amount of "sloppiness", since you will be verifying your guess later on.

$$\text{Ex: } T(n) = 3T(n/4) + \Theta(n^2)$$

We create a recursion tree for the recurrence: $T(n) = 3T(n/4) + cn^2$, 如图6所示

Assume that n is an exact power of 4 (another example of tolerable sloppiness)

Its height is $\log_4 n$ (it has $\log_4 n + 1$ levels $(0, 1, 2, \dots, \log_4 n)$)

The subproblem size for a node at depth i is $n/4^i$.




Figure 6: Recursion tree example

Each level has three times more nodes than the level above, and so the number of nodes at depth i is 3^i .

深度为 $i = 0, 1, 2, \dots, \log_4 n - 1$ 的 nodes 加起来是 $3^i c(n/4^i)^2 = (3/16)^i cn^2$.

最后一层, at depth $\log_4 n$, has $3^{\log_4 n} = n^{\log_4 3}$ nodes, each contributing cost $T(1)$, for a total cost of $n^{\log_4 3} T(1)$, which is $\Theta(n^{\log_4 3})$.

把所有层的加起来, 得到整个树的时间

$$\begin{aligned}
 T(n) &= cn^2 + \frac{3}{16}cn^2 + \left(\frac{3}{16}\right)^2 cn^2 + \dots + \left(\frac{3}{16}\right)^{\log_4 n - 1} cn^2 + \Theta(n^{\log_4 3}) \\
 &= \sum_{i=0}^{\log_4 n - 1} \left(\frac{3}{16}\right)^i cn^2 + \Theta(n^{\log_4 3}) \\
 &= \frac{(3/16)^{\log_4 n} - 1}{(3/16) - 1} cn^2 + \Theta(n^{\log_4 3})
 \end{aligned} \tag{2-3-1}$$

保留高阶项, 得到 $T(n) = O(n^2)$, 然后再用归纳法证明, 发现是正确的

2.4 Master method

Theorem: Let $a \geq 1$ and $b > 1$ be constants, let $f(n)$ be a asymptotically positive function, and let $T(n)$ be defined on the nonnegative integers by the recurrence

$$T(n) = aT(n/b) + f(n)$$

Where we interpret n/b to mean either $\lceil n/b \rceil$ or $\lfloor n/b \rfloor$. Then $T(n)$ can be bounded asymptotically as follows.

1. If $f(n) = O(n^{\log_b a - \epsilon})$ for some constant $\epsilon > 0$, then $T(n) = \Theta(n^{\log_b a})$
2. If $f(n) = n^{\log_b a}$, then $T(n) = \Theta(n^{\log_b a} \lg n)$
3. If $f(n) = n^{\log_b a + \epsilon}$ for some constant $\epsilon > 0$, and if $af(n/b) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n , then $T(n) = \Theta(f(n))$.

总结规律如下: Compare the function $f(n)$ with the function $n^{\log_b a}$

if $f(n)$ is polynomial smaller, case 1

if $f(n)$ is polynomial larger, case 3

It is important to realize that **these three cases do not cover all the possibilities for $f(n)$** . There is a gap between cases 1 and 2 when $f(n)$ is smaller than $n^{\log_b a}$ but not polynomially smaller. Similarly, there is a gap between cases 2 and 3 when $f(n)$ is larger than $n^{\log_b a}$ but not polynomially larger. If the function $f(n)$ falls into one of these gaps, or if the regularity condition in case 3 fails to hold, the master method cannot be used to solve the recurrence.

Ex: $T(n) = 3T(n/4) + n \lg n$,

we have $a = 3, b = 4, f(n) = n \lg n$,

$f(n)/(n^{\log_b a}) = n \lg n / (n^{\log_4 3}) = n^{1-\log_4 3} \times \lg n$

so $f(n)$ is polynomial larger, case 3, $T(n) = \Theta(n \lg n)$.

2.4.1 Proof of master method

树的深度是 $\log_b n$ 图示证明见7 不太严格的证明, 但是有助于理解




Figure 7: Recursion tree proof

case 3: 前面是呈几何级数递减的, 所以前面所有层的和 is dominated by the first member $f(n)$
 最后一层的和为 $a^{\log_b n} \times \Theta(1) = n^{\log_b a} \Theta(1) = \Theta(n^{\log_b a})$
 $T(n) = \Theta(f(n))$

所以比较 $f(n)$ and $(n^{\log_b a})$ 也就是比较最后一层与前面所有层的

case 1: 前面是呈几何级数递增的, 因为 $f(n)$ is polynomial smaller than $n^{\log_b a}$, 所以整个求和中, $n^{\log_b a}$ 占据主导

$T(n) = \Theta(n^{\log_b a})$

case 2: 由于最顶层和最底层的差不多, 在中间的, 从最顶层呢过到最底层变化又不大, 所以基本上都一个量级的, 所以总和是:

$$T(n) = (\log_b n + 1)f(n) = (\log_b n + 1) \times n^{\log_b a} = \Theta(f(n) \times \lg n)$$

3 分治法 divide and conquer

The divide-and-conquer design paradigm

1. Divide the problem (instance) into subproblems.
2. Conquer the subproblems by solving them recursively.
3. Combine subproblem solutions

3.1 Merge sort

1. Divide: Trivial.
2. Conquer: Recursively sort 2 subarrays.
3. Combine: Linear-time merge.

$$T(n) = 2 \times T(n/2) + \Theta(n) = \Theta(n \lg n)$$

3.2 Binary search

find x in **sorted array**

1. divide, compare x with middle
2. conquer : recurse in one subarray
3. combine: trivial

$$T(n) = \Theta(1) + 1 \times T(n/2) = \Theta(\lg n) \text{ master method}$$

Algorithm *BinarySearch*($L; x; first; last$)

Input: Array $L[first; last]$ and value x .

Output: 1 if $x \in L$, or i , $0 \leq i < n$ if $L[i] = x$

pseudo code

if $first > last$ then return -1

else {

$middle \leftarrow (first + last)/2$

if $L[middle] = x$ then return $middle$

else if $L[middle] < x$ then return *BinarySearch*($L; x; middle + 1; last$)

else return *BinarySearch*($L; x; first; middle - 1$)

}

3.3 Powering a number

given number x , integer $n \geq 0$, compute x^n

general method: $x \times x \times x \times x \times \dots \times x = x^n$, $T(n) = \Theta(n)$

Divide and conquer:

$$x^n = \begin{cases} x^{n/2} \times x^{n/2} & \text{if } n \text{ even} \\ x^{(n-1)/2} \times x^{(n-1)/2} \times x & \text{if } n \text{ odd} \end{cases}$$

$$T(n) = 1 \times T(n/2) + \Theta(1) = \Theta(\lg n)$$

$T(n/2)$ 表示算出平方根的时间, $\Theta(1)$ 表示将这个算出来的平方根平方的时间

3.4 Fibonacci numbers

$$F_0 = 0, F_1 = 1, F_2 = 1, F_3 = 2, \dots$$

通式: $F_n = F_{n-1} + F_{n-2}$

General method: $T(n) = \Omega(\phi^n)$ with $\phi = (1 + \sqrt{5})/2$ golden ratio

exponential time

bottom-up

Compute F_0, F_1, F_2, \dots in order

当我们要计算 F_n 的时候, 我们已经计算出了 F_{n-1} and F_{n-2} , 直接将这两个相加, 就得到 F_n , 将两个已知数相加所需要的时间为常数

所以只需要依次计算各个项就可以了

$$T(n) = n \times \Theta(1) = \Theta(n)$$

Naive recursive squaring

根据 Fibonacci 通项公式, 直接计算, and round to the nearest integer

但是浮点数的计算误差比较大

Recursive squaring

Théorème 3.1.

$$\begin{pmatrix} F_{n+1} & F_n \\ F_n & F_{n-1} \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}^n$$

prove the matrix equation by induction on n $T(n) = T(n/2) + \Theta(1)$ $T(n) = \Theta(\lg n)$

3.5 Matrix multiplication

input $A = a_{ij}, B = b_{ij}$

output $C = A \times B$

Standard algorithm

```
for i ← 1 to n
do for j ← 1 to n
do cij ← 0
for k ← 1 to n
do cij ← cij + aik \times bkj
```

直接计算, 三层嵌套循环, 时间复杂度是 $\Theta(n^3)$

divide and conquer algo:

直接分治法 $n \times n$ matrix = 2×2 matrix of $(n/2) \times (n/2)$ submatrices:

$$\mathbf{A} = \begin{bmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{bmatrix}, \mathbf{B} = \begin{bmatrix} \mathbf{B}_{1,1} & \mathbf{B}_{1,2} \\ \mathbf{B}_{2,1} & \mathbf{B}_{2,2} \end{bmatrix}, \mathbf{C} = \begin{bmatrix} \mathbf{C}_{1,1} & \mathbf{C}_{1,2} \\ \mathbf{C}_{2,1} & \mathbf{C}_{2,2} \end{bmatrix}$$

with

$$\mathbf{A}_{i,j}, \mathbf{B}_{i,j}, \mathbf{C}_{i,j} \in R^{2^{n-1} \times 2^{n-1}}$$

then

$$\mathbf{C}_{1,1} = \mathbf{A}_{1,1}\mathbf{B}_{1,1} + \mathbf{A}_{1,2}\mathbf{B}_{2,1}$$

$$\mathbf{C}_{1,2} = \mathbf{A}_{1,1}\mathbf{B}_{1,2} + \mathbf{A}_{1,2}\mathbf{B}_{2,2}$$

$$\mathbf{C}_{2,1} = \mathbf{A}_{2,1}\mathbf{B}_{1,1} + \mathbf{A}_{2,2}\mathbf{B}_{2,1}$$

$$\mathbf{C}_{2,2} = \mathbf{A}_{2,1}\mathbf{B}_{1,2} + \mathbf{A}_{2,2}\mathbf{B}_{2,2}$$

With this construction we have not reduced the number of multiplications. We still need 8 multiplications to calculate the $\mathbf{C}_{i,j}$ matrices, the same number of multiplications we need when using standard matrix multiplication.

Now comes the important part. We define new matrices

$$\mathbf{M}_1 := (\mathbf{A}_{1,1} + \mathbf{A}_{2,2})(\mathbf{B}_{1,1} + \mathbf{B}_{2,2})$$

$$\mathbf{M}_2 := (\mathbf{A}_{2,1} + \mathbf{A}_{2,2})\mathbf{B}_{1,1}$$

$$\mathbf{M}_3 := \mathbf{A}_{1,1}(\mathbf{B}_{1,2} - \mathbf{B}_{2,2})$$

$$\mathbf{M}_4 := \mathbf{A}_{2,2}(\mathbf{B}_{2,1} - \mathbf{B}_{1,1})$$

$$\mathbf{M}_5 := (\mathbf{A}_{1,1} + \mathbf{A}_{1,2})\mathbf{B}_{2,2}$$

$$\mathbf{M}_6 := (\mathbf{A}_{2,1} - \mathbf{A}_{1,1})(\mathbf{B}_{1,1} + \mathbf{B}_{1,2})$$

$$\mathbf{M}_7 := (\mathbf{A}_{1,2} - \mathbf{A}_{2,2})(\mathbf{B}_{2,1} + \mathbf{B}_{2,2})$$

only using 7 multiplications (one for each \mathbf{M}_k) instead of 8. We may now express the $\mathbf{C}_{i,j}$ in terms of \mathbf{M}_k , like this:

$$\mathbf{C}_{1,1} = \mathbf{M}_1 + \mathbf{M}_4 - \mathbf{M}_5 + \mathbf{M}_7$$

$$\mathbf{C}_{1,2} = \mathbf{M}_3 + \mathbf{M}_5$$

$$\mathbf{C}_{2,1} = \mathbf{M}_2 + \mathbf{M}_4$$

$$\mathbf{C}_{2,2} = \mathbf{M}_1 - \mathbf{M}_2 + \mathbf{M}_3 + \mathbf{M}_6$$

We iterate this division process n times (recursively) until the submatrices degenerate into numbers (elements of the ring R). The resulting product will be padded with zeroes just like A and B , and should be stripped of the corresponding rows and columns.

Strassen's algorithm divide and conquer

1. Divide: Partition A and B into $(n/2) \times (n/2)$ submatrices. Form terms to be multiplied using $+$ and $-$, time consumed: $\Theta(n^2)$
2. Conquer: Perform 7 multiplications(M_1, M_2, \dots, M_7) of $(n/2) \times (n/2)$ submatrices recursively. time consumed: $7 \times T(n/2)$

3. Combine: Form $C(r, s, t, u)$ using $+$ and $-$ on $(n/2) \times (n/2)$ submatrices. time consumed: $\Theta(n^2)$

时间复杂度

$$T(n) = 7 \times T(n/2) + \Theta(n^2) = \Theta(n^{\lg 7}) = \Theta(n^{2.807355})$$

当前最好的为 $n^{2.376}$ (理论上)

4 快排及随机算法

Tony Hoare 在 1962 年发明

-divide and conquer

-sorts "in place"(merge sort needs extra space, but qsort does not)

-very practical(with tuning)

In the worst case, it makes $O(n^2)$ comparisons, though this behavior is rare. Quicksort is often faster in practice than other $O(n \log n)$ algorithms.

Additionally, quicksort's sequential and localized memory references work well with a cache. Quicksort is a comparison sort and, in efficient implementations, is not a stable sort. Quicksort can be implemented with an in-place partitioning algorithm, so the entire sort can be done with only $O(\log n)$ additional space used by the stack during the recursion.

Algo

1 Divide: partition array into 2 sub arrays(见图8) around pivot(支点) such that elements in lower subarray $\leq x \leq$ elements in upper subarray

2 Conquer: recursively sort 2 subarrays




Figure 8: Quick sort partition

3 Combine: trivial(nothing to do for the combine)

key step in quicksort is partition step

快排在递归中也会进行 partition

key: linear-time $\Theta(n)$ partitioning subroutine

Partition(A,p,q) //A[p...q]

$x \leftarrow A[p]$ //pivot A[p]

$i \leftarrow p$

for $j \leftarrow p+1$ to q {

if $A[j] \leq x$ {

$i \leftarrow i+1$ //i加上1之后就到了大于x的那部分, 然后再将A[j] 交换到i 位置, loop invariant 保持 exchange A[i] with A[j]

}//end if

}//end for

exchange A[p] with A[i] // 将pivot 交换到中间, 完成partition 的任务

return i //end function

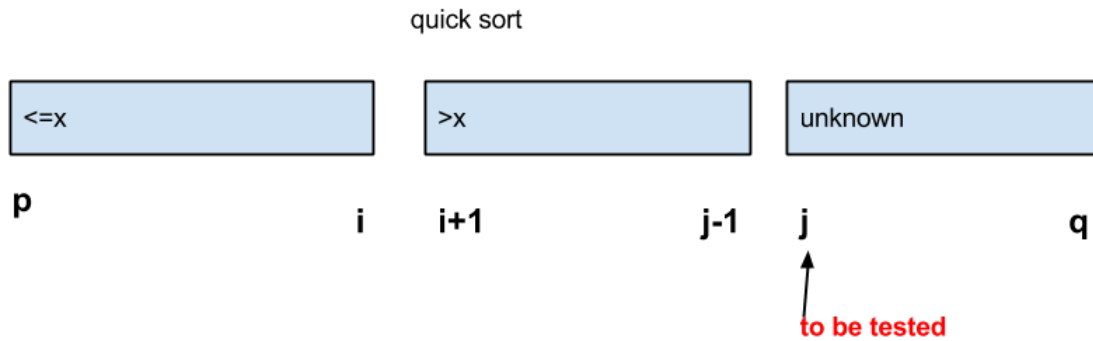


Figure 9: Quick sort loop invariant

At the beginning of each iteration of the loop, loop invariant

- if $k \in [p, i]$, then $A[k] \leq x$ and $A[p] = x$
- if $k \in [i + 1, j - 1]$, then $A[k] > x$
- if $k \in [j, q]$, then $A[k]$ unknown

running time $T(n)\Theta(n)$ (只需要遍历一遍) Ex: 28713564

在这个例子中, 最后一个元素被选为 pivot, 实际上 pivot 位于 p 或者 r 处都可以

```

Quicksort(A,p,q)
if p<q{
  r←Partition(A,p,q)
  Quicksort(A,p,r-1)
  Quicksort(A,r+1,q)
}
Initial call: Quicksort(A,1,n)

```

4.1 Analysis

当元素数目较少时,可以换用其他更快更直接的算法,这样可以避免再简单的情况下也进行递归。同时,这是一个 tail recursion(尾递归),所以可以使用 certain tail recursion optimizations

worst-case time

-input sorted or reverse sorted 元素都被分到了一边

so one side of partition has no elems, the other side has $n - 1$ elems

每次递归只减少了 1, 所以递归次数非常多

$$T(n) = \Theta(n) + T(0) + T(n-1) = \Theta(n) + \Theta(1) + T(n-1) = T(n-1) + \Theta(n)$$

式中 $\Theta(n)$ 表示 partition 所需要的时间

使用递归树得到时间复杂度为: $T(n) = \Theta(n^2)$

best-case analysis(intuition only!)

if we are really lucky, partition splits array $n/2 : n/2$

$T(n) = \Theta(n) + 2T(n/2) = \Theta(n \lg n)$ 和 merge sort 一样

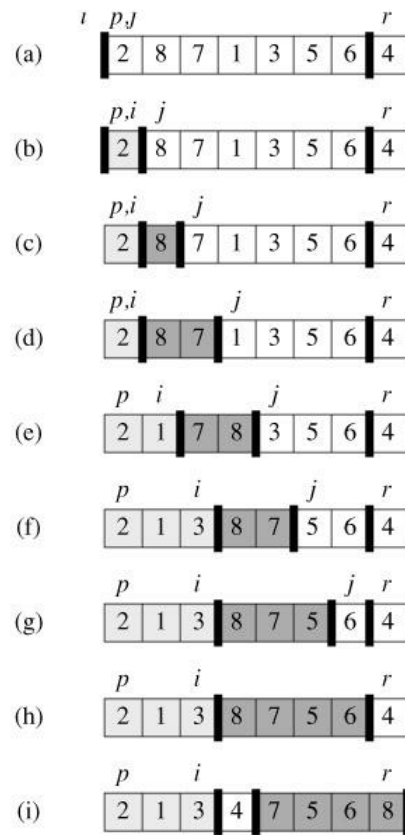


Figure 10: Quick sort example

Partition 1/10 : 9/10

$$T(n) = T(n/10) + T(9n/10) + \Theta(n)$$

用递归树得到

$$T(n) \leq (\log_{10/9} n * cn) + \Theta(n) = O(n \lg n)$$

$$T(n) \leq (\log_{10} n * cn) + \Theta(n) = \Omega(n \lg n)$$

$$T(n) = \Theta(n \lg n)$$

The reason is that any split of constant proportionality yields a recursion tree of depth $\Omega(\lg n)$ whenever the split has constant proportionality.

Suppose we alternate *lucky, unlucky, lucky...*

$$L(n) = 2U(n/2) + \Theta(n) \text{ lucky 的后面是两个 unlucky}$$

$$U(n) = L(n-1) + \Theta(n) \text{ unlucky}$$

Then

$$\begin{aligned} L(n) &= 2U(n/2) + \Theta(n) \\ &= 2(L(n/2-1) + \Theta(n/2)) + \Theta(n) \\ &= 2L(n/2-1) + \Theta(n) \\ &= \Theta(n \lg n) \text{ lucky} \end{aligned}$$

Suppose we alternate *unlucky, lucky, unlucky, lucky...*

$$\begin{aligned} U(n) &= L(n-1) + \Theta(n) \\ &= [2U((n-1)/2) + \Theta(n)] + \Theta(n) \\ &= 2U((n-1)/2) + \Theta(n) \\ &= 2U(n/2) + \Theta(n) \\ &= \Theta(n \lg n) \text{ lucky} \end{aligned}$$

So when we alternate lucky, unlucky, lucky..., we are lucky

Or we alternate unlucky, lucky, unlucky, lucky..., we are lucky too.

So how do we ensure that we are usually lucky?

Because if the input is already sorted or reverse sorted, we are going to be unlucky.

1. randomly arrange the elements

2. randomly choose the pivot

4.2 Random quicksort

pick the pivot randomly

选择好之后, 把这个选中的与 array 的第一个元素交换位置, 这样这个随机选择的 pivot 就到了 array 的第一位置, 然后再运行 Partition 函数

- 运行时间不取决于输入数据的顺序
- 对输入序列的分布不用做出假设
- 不存在特定的输入序列会引起 worst-case
- worst-case determined only by random number generator

4.3 Median-of-3 Pivot

For example, the median-of-3 pivot approach selects three candidate pivots and uses the median one. If the three pivots are chosen from the first, middle and last positions, then it is easy to see that for the already sorted array, this will produce an optimum result: each partition will be exactly half (\pm one element) of the problem and we will need exactly $\text{ceiling}(\log n)$ recursive calls.

4.4 Random method analysis

Random variable(随机变量) for running time assuming that random numbers are independent. I want to know where I pivoted. 设这个 pivot 的位置为随机变量 k

So for $k = 0, 1 \dots n-1$ let

$x_k = 1$ if partition generates a $k : n - k - 1$ split, (pivot 算一个数, 所以两者加起来是 $n-1$)

$x_k = 0$ otherwise

这样一个 partition, 就产生了 n 个 random variable, 其中只有一个是 1, 其余的都是 0.

例如如果产生的 partition 为 $5 : n - 6$, 那么只有 $x_5 = 1, x_i = 0 (i \in [[0, n-1]] \text{ and } i \neq 5)$

this type of random variable is called **indicator random variable**

the expected value of x_k :

$$E[x_k] = 0 * P(x_k = 0) + 1 * P(x_k = 1) = P(x_k = 1) = 1/n$$

因为每一个数 k 都有可能取到, 而且概率是一样的, 所以每个数的概率都是 $1/n$

$$T(n) = \begin{cases} T(0) + T(n-1) & \text{si } 0 : n-1 \text{ split} \\ T(1) + T(n-2) & \text{si } 1 : n-2 \text{ split} \\ \vdots \\ T(n-1) + T(0) & \text{si } n-1 : 0 \text{ split} \end{cases}$$

所以这就是 $T(n)$ 的递归, 但是这个递归很麻烦, 这里我们就可以看到 indicator random variable 的优美性, 我们将使用 indicator random variable 将这个递归 reduce to(规约到数学上)

$$T(n) = \sum_{k=0}^{n-1} x_k (T(k) + T(n-k-1) + \Theta(n))$$

$T(n)$ 的期望

$$\begin{aligned}
 E[T(n)] &= E\left[\sum_{k=0}^{n-1} x_k (T(k) + T(n-k-1) + \Theta(n))\right] \\
 &= \sum_{k=0}^{n-1} E[x_k (T(k) + T(n-k-1) + \Theta(n))] \\
 &\quad x_k \text{ 随机变量独立于任何其他 partitions, 也就是说 } x_k \text{ 区别于} \\
 &\quad \text{其他递归调用, 所以积的期望等于期望的积} \tag{4-4-1} \\
 &= \sum_{k=0}^{n-1} E[x_k] * E[(T(k) + T(n-k-1) + \Theta(n))] \\
 &= \frac{1}{n} \sum_{k=0}^{n-1} E[T(k)] + \frac{1}{n} \sum_{k=0}^{n-1} E[T(n-k-1)] + \frac{1}{n} \sum_{k=0}^{n-1} \Theta(n) \\
 &= \frac{2}{n} \sum_{k=0}^{n-1} E[T(k)] + \Theta(n)
 \end{aligned}$$

x_k 是对 n 进行 partition 时产生的一个 indicator random variable 也就是说 x_k 是和 $T(n)$ 是相关的, 而后面的 $T(k)$ 与 $T(n-k-1)$ 是在有了 x_k 之后, 也就是产生了一个 $k, n-k-1$ 的分割后, 对产生的两个新的分割求时间时, 又会产生新的 indicator: x'_k

所以 x_k 与 $T(k), T(n-k-1)$ 是相互独立的, 可以运用概率的乘法原则

Absorb $k=0, 1$ terms into $\Theta(n)$ for tech convenience, (两个常量加进去不影响 $\Theta(n)$)

$$E[T(n)] = \frac{2}{n} \sum_{k=2}^{n-1} E[T(k)] + \Theta(n)$$

Use fact

$$\sum_{k=2}^{n-1} k \lg k \leq \frac{1}{2} n^2 \lg n - \frac{1}{8} n^2$$

how to
prove?

To prove $E[T(n)] \leq a * n \lg n$ for constant $a > 0$ with the substitution method

$$E[T(n)] \leq an \lg n - bn, (a, b > 0)$$

.....

所以我们得到 $T(n)$ 的期望值是 $\Theta(n \lg n)$

so the running time of randomized quicksort is $\Theta(n \lg n)$

The version of PARTITION given in this chapter is not the original partitioning algorithm.

Here is the original partition algorithm, which is due to T.Hoare:

HOARE-PARTITION(A, p, r)

```

1  x ← A[p]
2  i ← p - 1
3  j ← r + 1
4  while TRUE
5      do repeat j ← j - 1
6          until A[j] ≤ x
7      repeat i ← i + 1

```

```

8         until A[i] >= x
9         if i < j
10            then exchange A[i] with A[j]
11            else return j

```

Every element of $A[p..j]$ is less than or equal to every element of $A[j+1..r]$ when HOARE-PARTITION terminates.

The HOARE-PARTITION procedure always places the pivot value (originally in $A[p]$) into one of the two partitions $A[p..j]$ and $A[j+1..r]$.

4.5 Exercise

[Exercises 7.4-5] The running time of quicksort can be improved in practice by taking advantage of the fast running time of insertion sort when its input is "nearly" sorted. When quicksort is called on a subarray with fewer than k elements, let it simply return without sorting the subarray. After the top-level call to quicksort returns, run insertion sort on the entire array to finish the sorting process. Argue that this sorting algorithm runs in $O(nk + n \lg(n/k))$ expected time. How should k be picked, both in theory and in practice?

[Problems 7-6: Fuzzy sorting of intervals] Consider a sorting problem in which the numbers are not known exactly. Instead, for each number, we know an interval on the real line to which it belongs. That is, we are given n closed intervals of the form $[a_i, b_i]$, where $a_i \leq b_i$. The goal is to fuzzy-sort these intervals, i.e., to produce a permutation i_1, i_2, \dots, i_n of the intervals such that for $j = 1, 2, \dots, n$ there exist $c_j \in [a_{i_j}, b_{i_j}]$, satisfying $c_1 \leq c_2 \leq \dots \leq c_n$.

1. Design an algorithm for fuzzy-sorting n intervals. Your algorithm should have the general structure of an algorithm that quicksorts the left endpoints (the a_i values), but it should take advantage of overlapping intervals to improve the running time. (As the intervals overlap more and more, the problem of fuzzy-sorting the intervals gets easier and easier. Your algorithm should take advantage of such overlapping, to the extent that it exists.)
2. Argue that your algorithm runs in expected time $\Theta(n \lg n)$ in general, but runs in expected time $\Theta(n)$ when all of the intervals overlap (i.e., when there exists a value x such that $x \in [a_i, b_i]$ for all i). Your algorithm should not be checking for this case explicitly; rather, its performance should naturally improve as the amount of overlap increases.

5 Heapsort

Implementation of heap: tree or array

Max-heap: parent \geq children, heapsort

Min-heap: parent \leq children, priority queue

下面我们以 Max-heap 进行讲解

Array n elements, leaves: $\text{floor}(n/2) + 1, \text{floor}(n/2) + 2, \text{floor}(n/2) + 3, \dots, n$

The (binary) heap data structure is an array object that we can view as a nearly complete binary tree

```

PARENT(i)
    return floor(i/2)
LEFT(i)
    return 2i
RIGHT(i)
    return 2i + 1

```




Figure 11: An example of heap

The height of a node: the length of the longest downward path to a leaf from that node

The depth of a node: the length of the path to its root

Root node has depth 0, leaf node has height 0

5.1 Maintaining the heap property

Its inputs are an array A and an index i into the array

$A[i]$ "float down"

```

Max-heapify(A, i)
    l = left(i)
    r = right(i)
    if l <= A.heap-size and A[l] > A[i]
        largest = l
    else largest = i
    if r <= A.heap-size and A[r] > A[i]
        largest = r
    if largest != i
        exchange A[i] with A[largest]
        Max-heapify(A, largest)

```

max heapify: correct a single violation of the heap property in a sub tree's root

运行 **Max-heapify**(A, i) 的前提条件 **precondition**: assume that the trees rooted at $\text{left}(i)$ and $\text{right}(i)$ are max-heaps

图12中的第二个节点 violates the max heap property, 所以我们可以调用 $\text{max-heapify}(A, 2)$, 修复好这个之后, 再进行检查直到不出现 violation 为止




Figure 12: An example of violation of heap property

The children's subtrees each have size at most $2n/3$ —the worst case occurs when the bottom level of the tree is exactly half full—and therefore we can describe the running time of MAX-HEAPIFY by the recurrence

$$T(n) \leq T(2n/3) + \Theta(1)$$

master theorem $\Rightarrow T(n) = O(\lg n)$

add a new item : place the new item at the last position(向左向下走知道不能继续前进), 然后进行修复, compare it with its parent, if it is larger than its parent, exchange them 向上继续进行这个比较直到不违反 heap 的性质, 也就是向上修复

delete the root : 将最后一个节点的值赋给 root, 然后向下进行修复

not understood

5.2 Building a heap

build-max heap: produce a max heap from an unordered array

The elements in the subarray $A[(\text{floor}(n/2) + 1) \dots n]$ are all leaves of the tree 而 leaf 已经是 max-heap 了, 因为 leaf 没有 children

The procedure BUILD-MAX-HEAP goes through the remaining nodes of the tree and runs MAX-HEAPIFY on each one

Build-Max-heap(A)

A.heap-size = A.length

for i = floor(A.length / 2) downto 1

Max-heapify(A, i) //每次调用max-heapify 之前, 都是满足max-heapify的precondition 的

Each call to MAX-HEAPIFY costs $O(\lg n)$ time, and BUILD-MAX-HEAP makes $O(n)$ such calls. Thus, the running time is $O(n \lg n)$. This upper bound, though correct, is not asymptotically tight.

但是我们可以发现:

max-heapify takes $O(1)$ for nodes that are one level above the leaves

and in general $O(l)$ time for nodes that are l levels above the leaves

how to prove

an n -element heap has height $\text{floor}(\lg n)$ and at most $\text{ceil}(n/(2^{h+1}))$ nodes of any height h need proof

$$\sum_{h=0}^{\lfloor \lg n \rfloor} \lceil \frac{n}{2^{h+1}} \rceil O(h) = O(n \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{h}{2^h})$$

$$\sum_{h=0}^{\infty} \frac{h}{2^h} = \frac{1/2}{(1 - 1/2)^2} = 2$$

Thus, we can bound the running time of BUILD-MAX-HEAP as

$$O(n \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{h}{2^h}) = O(n \sum_{h=0}^{\infty} \frac{h}{2^h}) = O(n)$$

Hence, we can build a max-heap from an unordered array in linear time

5.3 The heapsort algorithm

build-max-heap from unordered array

find max element $A[1]$

swap elements $A[n]$ with $A[1]$, now max element is at the end of the array

Discard node n from heap,

New root may violate max heap property, fix it

Heapsort(A)

 Build-Max-heap(A)

 for $i = A.\text{length}$ downto 2

 exchange $A[1]$ with $A[i]$

$A.\text{heap-size} = A.\text{heap-size} - 1$

 Max-heapify(A, 1)

把 $A[1]$ 也就是最大值挪到最后一个位置, 同时将 heap 的大小减一, 这样最大值就从 heap 中取出来了并被放在了正确的位置上, 然后再修复这个 heap; 然后再最大值重复这个操作

time $O(n \lg n)$

in place sort

5.4 Priority queues

max-priority queue, min-priority queue

we can use a heap to implement a priority queue

HEAP-MAXIMUM(A)

 return $A[1]$

heap-maximum: $\Theta(1)$ time

HEAP-extract-MAX(A)

 if $A.\text{heap-size} < 1$

```

        error "heap underflow"
    max = A[1]
    A[1] = A[A.heap-size]
    A.heap-size = A.heap-size - 1
    MAX-HEAPIFY(A,1)
    return max

```

heap-extract-max: $O(\lg n)$ time

```

HEAP-INCREASE-KEY(A,i,key)
    if key < A[i]
        error "new key is smaller than current key"
    A[i] = key
    while i > 1 and A[PARENT(i)] < A[i]
        exchange A[i] with A[PARENT(i)]
        i = PARENT(i)

```

HEAP-INCREASE-KEY: $O(\lg n)$ time

```

MAX-HEAP-INSERT(A,key)
    A.heap-size = A.heap-size + 1
    A[A.heap-size] = -\infty
    HEAP-INCREASE-KEY(A,A.heap-size,key)

```

MAX-HEAP-INSERT: $O(\lg n)$ time

6 线性时间排序

- quicksort $\Theta(n \lg n)$ randomized
- heapsort $\Theta(n \lg n)$ 堆排序
- merge sort $\Theta(n \lg n)$
- insertion sort $\Theta(n^2)$

can we do better than $\Theta(n \lg n)$?

comparison sorting model

only use comparisons to determine relative order of elements

6.1 decision tree model 决策树

Ex: sort $\langle a_1, a_2, a_3 \rangle$

In general $\langle a_1, a_2, \dots, a_n \rangle$

- each internal node labeled i, j means compare a_i vs. a_j

- left subtree $a_i \leq a_j$
- right subtree $a_i > a_j$
- each leaf gives a permutation of the array

图13是一个决策树的例子




Figure 13: Decision tree example

A lower bound for the worst case

The length of the longest path from the root of a decision tree to any of its reachable leaves represents the worst-case number of comparisons. Consequently, the worst-case number of comparisons for a given comparison sort algorithm equals the height of its decision tree. A lower bound on the heights of all decision trees in which each permutation appears as a reachable leaf is therefore a lower bound on the running time of any comparison sort algorithm. The following theorem establishes such a lower bound.

[Theorem 8.1] Any comparison sort algorithm requires $\Omega(n \lg n)$ comparisons in the worst case. Proof: From the preceding discussion, it suffices to determine the height of a decision tree in which each permutation appears as a reachable leaf.

Consider a decision tree of height h with l reachable leaves corresponding to a comparison sort on n elements. Because each of the $n!$ permutations of the input appears as some leaf, we have $n! \leq l$. (至少含有 $n!$ 个 leaves) Since a binary tree of height h has no more than 2^h leaves, we have $n! \leq l \leq 2^h$, which, by taking logarithms, implies $h \geq \lg(n!) = \Omega(n \lg n)$ (Stirling's approximation $n! \simeq (n/e)^n$)

6.2 Counting sort 计数排序

Element $\in [0, k]$ for some integer k .

When $k = O(n)$, the sort runs in $\Theta(n)$ time.

The **basic idea** of counting sort is to **determine, for each input element x , the number of elements less than x** . This information can be used to place element x directly into its position in the output array.

For example, if there are 17 elements less than x , then x belongs in output position 18. This scheme must be modified slightly to handle the situation in which several elements have the same value, since we don't want to put them all in the same position.

input array $A[1, \dots, n]$, and thus $\text{length}[A] = n$.
 the array $B[1, \dots, n]$ holds the sorted output,
 the array $C[0, \dots, k]$ provides temporary working storage.

COUNTING-SORT(A, B, k)

```

1  for i ← 0 to k
2      do C[i] ← 0
3  for j ← 1 to length[A]
4      do C[A[j]] ← C[A[j]] + 1
5  // C[i] now contains the number of elements equal to i.
6  for i ← 1 to k
7      do C[i] ← C[i] + C[i - 1]
8  // C[i] now contains the number of elements leqi
9  for j ← length[A] downto 1
10     do B[C[A[j]]] ← A[j]
11     C[A[j]] ← C[A[j]] - 1
```

An important property of counting sort is that it is **stable**.

Analysis

line 1 – $2\Theta(k)$

line 3 – $4\Theta(n)$

line 6 – $7\Theta(k)$

line 9 – $11\Theta(n)$

Overall time $\Theta(n + k)$

在实际上, 当 $k\Theta(n)$ 时, 我们才采用 counting sort

[Exercises 8.2-3] Suppose that the for loop header in line 9 of the COUNTING-SORT procedure is rewritten as

```
9  for j ← 1 to length[A]
```

Show that the algorithm still works properly. Is the modified algorithm stable?

凭直觉来看, 不是 stable

6.3 Radix sort 基数排序

LSD(Least Significant Digit first) 的基数排序适用于位数小的数列, 如果位数多的话, 使用 MSD 的效率会比较好.

MSD(Most Significant Digit first) 的方式与 LSD 相反, 是由高位数为基底开始进行分配, 但在分配之后并不马上合并回一个数组中, 而是在每个”桶子”中建立”子桶”, 将每个桶子中的数值按照下一数位的值分配到”子桶”中. 在进行完最低位数的分配后再合并回单一的数组中.

这里讲到的是 LSD

6.3.1 说明

其原理在于对于待排序的数据, 整体权重未知的情况下, 先按权重小的因子排序, 然后按权重大的因子排序.

例如比较时间, 先按日排序, 再按月排序, 最后按年排序, 仅需排序三次.

但是如果先排序高位就没这么简单了.

基数排序源于老式穿孔机, 排序器每次只能看到一个列 (这就是一个整体权重未知的例子), 很多教科书上的基数排序都是对数值排序 (这样意义不大), 数值的大小是已知的, 与老式穿孔机不同. 将数值按位拆分再排序, 是无聊并自找麻烦的事. 算法的目的是找到最佳解决问题的方案, 而不是把简单的事搞的更复杂.

基数排序更适合用于对时间, 字符串等这些整体权值未知的数据进行排序. 这时候基数排序的思想才能体现出来, 例如字符串, 如果从高位 (第一位) 往后排就很麻烦. 而反过来, 先对影响力较小, 的低位 (最后一位) 进行排序就非常简单了. 这时候基数排序的思想就能体现出来.

d 位数字, 每一位数字可以取 k 个不同的数, 例如 0 到 9 十个数, 那么 $k = 10$

RADIX-SORT(A, d)

```
for i from 1 to d:
    do use a stable sort to sort array A on digit i
```

6.3.2 Analysis

time: $\Theta(d(n + k))$

证明: correctness

induct on digit position t

assume by induction 前 $t - 1$ 位已经排好序, 我们需要对第 t 位进行排序

1, if two elements have same t^{th} digit, stability \Rightarrow same order \Rightarrow sorted order

2, if different t^{th} digit \Rightarrow sorted order

—use counting sort digits

—suppose we have n integers each b bits ($range = [0, 2^b - 1]$, non-negative)

—split into b/r "digits" each r bits (将 r 个 bits 组合成一个大的"bit")

time $O(b/r * (n + k)) = O(b/r * (n + 2^r))$

对 r 求导, 令其为零求得此时的 r 值 $r = \lg n$

得到 $O(bn/\lg n)$

if number in range $0, 2^b - 1$ then $time = O(dn)$

if $d = O(1)$ then $time = O(n)$

而且只要 d 小于 $\lg n$, 就可以击败 comparison sort

但是在实际上, counting sort is not very good on a cache, in practice, radix sort is not that fast, unless your numbers are really small.

但是在理论上, 这个算法很优美

Finally, if you have arbitrary integers, that are one word length long, and you can manipulate a word in constant time.

Then the best algorithm we known for sorting runs $n * \sqrt{\lg \lg n}$, and this is a randomized algorithm and very complicated

有另外一个算法 $n * \lg(\lg n)$ worst case, 这篇论文应该可以看懂

7 Order statistics

given n elements

find k th smallest element

naive algorithm: sort and return $A[k]$

$k = 1$ minimum

$k = n$ maximum

median 中位数 $k = \text{floor}((n+1)/2)$ or $\text{ceil}((n+1)/2)$

randomized divide and conquer

random-select

$\Theta(n)$ expected running time

$\Theta(n^2)$ worst-case 差不多是 $1/n^n$ 的概率

intuition for analysis

(we assume 所有数都不等)

lucky case: $1/10, 9/10$

$T(n) \leq T(9/10n) + \Theta(n)$ 假设第 i 小的位于 $9/10$ 部分

unlucky case: $0, n-1$

$T(n) = T(n-1) + \Theta(n)$

indicator random variable

let $T(n)$ be the random variable for running time of random-select

define indicator random variable $x_k, k \in [0, n-1]$

$x_k = 1$ if partition generates an k to $n-k-1$ splits

$x_k = 0$ otherwise

substitution method, $E[T(n)] \leq cn$

$3/8n^2$ induction to get $3/8$

worst-case linear time order statistics

Why do we group in 5 numbers, not 3 or 7 or others?

3 是不行的, 7 也行, 但是对于效率没有什么提升

8 Hashing

symbol-table problem in compiler

table S holding n symbols

dictionary operations

$\text{inset}(S, x)$ x is an element

$\text{delete}(S, x)$

$\text{search}(S, k)$ for a given key

search time $\Theta(n)$ in the worst case

search time $\Theta(1)$ the expected time

8.1 Direct address tables

Direct addressing is applicable when we can afford to allocate an array that has one position for every possible key.

```
Direct-address-search(T, k)
    return T[k]
```

```
Direct-address-insert(T, x)
    T[key[x]] = x
```

```
Direct-address-delete(T, x)
    T[key[x]] = nil
```

8.2 Hash tables

With direct addressing, an element with key k is stored in slot k .

With hashing, this element is stored in slot $h(k)$; that is we use a **hash function** h to compute the slot from the key k .

Here h maps the universe U of keys into the slots of a **hash table** $T[0, \dots, m-1]$

$h : U \rightarrow 0, 1, \dots, m-1$

collision: two keys hash to the same slot.

Solution: make h appear random thus avoiding collisions or at least minimizing their number.

Since $|U| > m$, 肯定会有 collisions

Collision resolution by chaining

We put all elements that hash to the same slot in a linked list.

Slot j contains a pointer to the head of list of all stored elements that hash to j .

通过这种方式解决 collision 的 dictionary operations 很容易实现

```
chain-hash-insert(T, x)
    insert x at the head of list T[h(key[x])]
```

```
chain-hash-search(T, k)
    search for an element with key k in list T[h(k)]
```

```
chain-hash-delete(T, x)
    delete x from the list T[h(key[x])]
```

Analysis

Insert: worst-case $O(1)$

插入很快, 是因为我们假设 hash table 之前不含有 x ; 当然我们可以在 insert 之前先 search 一遍

Search: proportional to the length of linked list

Delete: $O(1)$ if the list is doubly linked

same as search if the list is singly linked

How well does hashing with chaining perform? In particular, how long does it take to search for an element with a given key?

Def.: the **load factor** of a hash table with n keys and m slots is

$\alpha = n/m$ = average keys in a slot = average number of elements stored in a chain

worst-case

集合 S 中的所有键经过 hash 函数生成的 hash 值是一样的, 因此 all elements go to the same slot
the access time is $\Theta(n)$ + the time to compute the hash function.

average-case

assumption of **simple uniform hashing**

—any given element is equally likely to hash into any of the m slots, independently of where any other element has hashed to.

For $i \in [0, m-1]$, let us denote n_i the length of the list $T[i]$

So that $n = n_0 + n_1 + \dots + n_{m-1}$

And the average value of n_i is $E[n_i] = \alpha = n/m$

expected unsuccessful search time $\Theta(1 + n/m)$ 其中 1 是指进行 hash 操作的时间

expected successful search time = $\Theta(1 + n/m)$

Page 227: the situation for a successful search is slightly different, since each list is not equally likely to be searched. Instead, the probability that a list is searched is proportional to the number of elements it contains. Nonetheless, the expected search time is still $\Theta(1 + \alpha)$

Proof: the number of elements examined during a successful search for an element x is 1 plus the number of elements that appear before x in x 's list. Elements before x in the list were all inserted after x was inserted.

Let x_i denote the i^{th} element inserted into the table, and let $k_i = \text{key}[x_i]$

For keys k_i and k_j (k_j 注意细微差别),

we define indicator random variable $X_{ij} = I\{h(k_i) = h(k_j)\} = h(k_j)$

Under the assumption of simple uniform hashing, we have $\Pr\{h(k_i) = h(k_j)\} = 1/m$

So $E[X_{ij}] = 1/m$

For x_i , the number of elements examined in a successful search is $(1 + \sum_{j=i+1}^n X_{ij})$, 其中的 1 是 x_i 本身, 他自己也要被检查

Thus, the expected number of elements examined in a successful search is:

$$E\left[\frac{1}{n} \sum_{i=1}^n \left(1 + \sum_{j=i+1}^n X_{ij}\right)\right] = 1 + \frac{n-1}{2m} = 1 + \frac{\alpha}{2} - \frac{\alpha}{2n}$$

Thus the total time required for a successful search (including the time for computing the hash function) is $\Theta(2 + \alpha/2 - \alpha/2n) = \Theta(1 + \alpha)$

if that is to say When the number of hash-table slots is at least proportional to the number of elements in the table, that is to say $n\Theta(m)$ and $\alpha\Theta(1)$. Thus searching takes constant time on average.

Since insertion takes $O(1)$ worst-case time and deletion takes $O(1)$ worst-case time when the lists are doubly linked, all dictionary operations can be supported in $O(1)$ time on average.

为什么是成比例的, 搜索哪个 list 不是通过计算 key 的 hash 值来确定的吗?

how to compute?

8.3 Hash functions

A good hash function satisfies (approximately) the assumption of simple uniform hashing.

但是我们很少能够知道 keys 的 distribution, 所以这个条件很难满足.

如果我们知道 distribution, 例如, if the keys are known to be random real numbers k independently and uniformly distributed in the range $0 \leq k < 1$, the hash function

$h(k) = \text{floor}(km)$ 满足 simple uniform hashing 的条件.

In practice, heuristic(启发式的) techniques can often be used to create a good hash function.

Interpreting keys into natural numbers:

Most hash functions assume that the universe of keys is the set $N = 0, 1, 2, \dots$ of natural numbers 因此, 当 keys 不是自然数时, 需要将其转化为自然数.

例如, a character string \rightarrow natural number **in suitable radix notation**

For example: the string pt, (112, 116) in ASCII character set, then expressed as a radix-128 integer, pt becomes $112 * 128 + 116 = 14452$

to be
checked

8.3.1 Division method

$$h(k) = k \bmod m$$

When using the division method, we usually avoid certain values of m .

一个数对 2^r 取余, 结果就是这个数的二进制表示的最后 r 位数字

同样的道理, 一个数对 10^r 取余, 结果就是这个数的 10 进制表示的最后 r 位数字

所以选择一个质数来作为 m 比较好, 这个质数一定不要接近 2 或者 10 的某个指数

8.3.2 Multiplication method

$$h(k) = \text{floor}(m(kA \bmod 1))$$

where $x \bmod 1$ means the fractional part of x , that is $x - \text{floor}(x)$

A is constant $\in (0, 1)$

An advantage of the multiplication method is that the value of m is not critical. We typically choose it to be a power of 2 ($m = 2^p$ for some integer p). 这样 multiplication method 就很容易在计算机上实现

具体执行方法

word size of the machine: w bits

k fits into a single word

A is a fraction of the form $s/2^w$, where s is an integer in the range $0 < s < 2^w$.

we first multiply k by the w -bit integer $s = A \cdot 2^w$. The result is a $2w$ -bit value $r_1 * 2^w + r_0$, where r_1 is the high-order word of the product and r_0 is the low-order word of the product.

The desired p -bit hash value consists of the p most significant bits of r_0 .

Although this method works with any value of the constant A , it works better with some values than with others.

$$A = (\sqrt{5} - 1)/2 = 0.6180339887$$

Example: suppose we have $k = 123456$, $p = 14$, $m = 2^{14} = 16384$, and $w = 32$. Adapting Knuth's suggestion, we choose A to be the fraction of the form $s/2^{32}$ that is closest to , so that




Figure 14: Multiplication method for Hash

$A = 2654435769/232$.

Then $k \times s = 327706022297664 = (76300 \times 2^{32}) + 17612864$,

and so $r_1 = 76300$ and $r_0 = 17612864$

The 14 most significant bits of r_0 yield the value $h(k) = 67$.

17612864 的二进制表示为: 1000,0110,0110,0000,0010,0000,0

$m = 2^p$ and word size of machine is w -bit,

k fits into a single word

$h(k) = (A * k \bmod 2^w) \text{rsh}(w - r)$

A an odd integer in the range $(2^{w-r} - 1, 2^w)$, don't pick A too close to 2^{m-1} or 2^w

rsh right shift

14bits 为什么只有 67, 这么小??? 不知道这个是怎么算出来的

这个是 MIT 视频教程中使用的方法

8.3.3 Collision resolution by open addressing

all elements are stored in the hash table itself.

– no storage for links

进行多次 hash 操作直到找到一个空的 slot

probe sequence should be a permutation

table may full up $n \leq m$

deletion is difficult

To perform insertion using open addressing, we successively examine, or probe, the hash table until we find an empty slot in which to put the key. Instead of being fixed in the order $0, 1, \dots, m-1$ (which requires $\Theta(n)$ search time), the sequence of positions probed depends upon the key being inserted. To determine which slots to probe, we extend the hash function to include the probe number (starting from 0) as a second input. Thus, the hash function becomes

$h : U \times 0, 1, \dots, m-1 \rightarrow 0, 1, \dots, m-1$.

With open addressing, we require that for every key k , the probe sequence

$h(k, 0), h(k, 1), \dots, h(k, m-1)$

be a permutation of $0, 1, \dots, m-1$, so that every hash-table position is eventually considered as a slot for a new key as the table fills up. In the following pseudocode, we assume that the elements in the hash table T are keys with no satellite information; the key k is identical to the element containing key k . Each slot contains either a key or *NIL* (if the slot is empty).

这一段没有怎么看懂. 书上的 237 页

HASH-INSERT(T, k)

```

1  i ← 0
2  repeat j ← h(k, i)
3      if T[j] = NIL
4          then T[j] ← k
5          return j
6      else i ← i + 1
7  until i = m
8  error "hash table overflow"

```

搜索用到的 sequence 和 insert 用到的 sequence 一样

HASH-SEARCH(T, k)

```

1  i ← 0
2  repeat j ← h(k, i)
3      if T[j] = k
4          then return j
5      i ← i + 1
6  until T[j] = NIL or i = m
7  return NIL

```

Deletion from an open-address hash table is difficult.

对于 delete, 我们不能简单的将其标记为 nil, 这样在 search 的时候就会出现问題.

例如, 插入 k_1 的时候, i 位置是有值的, 所以 k_1 会在 i 之后继续 probe; 然后 i 位置的值被删除, 被标记成了 *nil*, 这样在查找 k_1 的时候, 当试探到 i 位置的时候, 发现是 *nil*, 达到 search 中止的条件, 返回 *nil*, 但是 k_1 确实存在于 hash table.

所以, 删除不能简单地标记为 *nil*, 我们可以标记为 DELETED, 这样 insert 的代码需要稍作修改, search 的代码不用动

但是又产生了另外一个问题, search times are no longer dependent on the load factor , and for this reason chaining is more commonly selected as a collision resolution technique when keys must be deleted.

probe strategy

1, linear probing:

Given an ordinary hash function h' , which we refer to as an auxiliary hash function.

$$h(k, i) = (h'(k) + i) \bmod m$$

Because the initial probe determines the entire probe sequence, there are only m distinct probe sequences.

这个方案不好, 如果碰到 primary clustering 集群现象, 也就是在一个空的 slot 前面全是 full slots, 所以每次探查的效率不高

$$2, \text{quadratic probing } h(k, i) = (h'(k) + c_1 * i + c_2 * i^2) \bmod m$$

better than linear probing

but to make full use of hash table, c_1, c_2 and m are constrained.

milder 温和的 form of clustering, secondary clustering

与 linear probing 一样, 只有 m 个不同 probe sequence

书上没有展
开讲解

3, double hashing

$$h(k, i) = (h_1(k) + i * h_2(k)) \bmod m$$

one way pick $m = 2^r$ and $h_2(k)$ always returns an odd 奇数 number

another way is to let m be prime and h_2 always returns a positive integer less than m

For example:

$$h_1(k) = k \bmod m$$

$$h_2(k) = 1 + (k \bmod m') \text{ where } m' \text{ is chosen slightly less than } m \text{ (say } m - 1)$$

$\Theta(m^2)$ probe sequences, rather than $\Theta(m)$

未理解

For uniform hashing, there are $m!$ permutations, so $m!$ probe sequences.

Analysis of open addressing hashing

With open addressing, we have at most one element per slot, and thus $n \leq m$, which implies $\alpha \leq 1$

Theorem: Given an open-address hash table with load factor $\alpha < 1$, the expected number of probes in an *unsuccessful search* is at most $1/(1 - \alpha)$, assuming uniform hashing

Proof: In an unsuccessful search, every probe but the last accesses an occupied slot that does not contain the desired key, and the last slot probed is empty.

Let us define the random variable X to be the number of probes made in an unsuccessful search, and let us also define the event A_i , for $i \in [1, \infty]$ to be the event that there is i^{th} probe and it is an occupied slot. Then **the event $\{X \geq i\}$ is the intersection of events**

$$A_1 \cap A_2 \cap \dots \cap A_{i-1}$$

$$\begin{aligned} Pr\{X \geq i\} &= Pr\{A_1 \cap A_2 \cap \dots \cap A_{i-1}\} \\ &= Pr\{A_1\} * Pr\{A_2|A_1\} * Pr\{A_3|A_1 \cap A_2\} * \dots * Pr\{A_{i-1}|A_1 \cap A_2 \cap \dots \cap A_{i-2}\} \end{aligned}$$

Since there are n elements and m slots, $Pr\{A_1\} = n/m$.

$Pr\{A_2|A_1\}$ 由于 A_1 是确定发生的, 所以第一个 slot 已经被占用, 所以还剩下 $m - 1$ 个 slots, 同时还剩下 $n - 1$ 个 elements, 所以 $Pr\{A_2|A_1\} = (n - 1)/(m - 1)$

所以

$$\begin{aligned} Pr\{X \geq i\} &= n/m * (n - 1)/(m - 1) * (n - 2)/(m - 2) * \dots * (n - i + 2)/(m - i + 2) \\ &\leq (n/m)^{i-1} \\ &= (\alpha)^{i-1} \end{aligned}$$

$$E[X] = \sum_{i=1}^{\infty} i * Pr\{X = i\} = \sum_{i=1}^{\infty} Pr\{X \geq i\} \leq \sum_{i=1}^{\infty} (\alpha)^{i-1} = \sum_{i=0}^{\infty} (\alpha)^i = 1/(1 - \alpha)$$

table 50% full, 2 次

table 90% full, 10 次

Theorem: Given an open-address hash table with load factor $\alpha < 1$, the expected number of probes in an *successful search* is at most

$1/\alpha * \ln 1/(1 - \alpha)$ that is $1/\alpha * \ln 1/(1 - \alpha)$, assuming uniform hashing and assuming that each key in the table is equally likely to be searched for.

Proof: A search for a key k follows the same probe sequence as was followed when the element with key k was inserted. If k was the $(i + 1)^{\text{st}}$ key inserted into the hash table, the expected number of

probes made in a search for k is at most $1/(1 - i/m) = m/(m - i)$

(k 是 $i + 1$ 个插入的, 说明 hash table 中已经含有 i 个 element, 所以此时的 load factor i/m)

Averaging over all n keys in the hash table gives us the average number of probes in a successful search:

$$\begin{aligned}\frac{1}{n} \sum_{i=0}^{n-1} \frac{m}{m-i} &= \frac{m}{n} \sum_{i=0}^{n-1} \frac{1}{m-i} \\ &= \frac{1}{\alpha} (H_m - H_{m-n})\end{aligned}$$

where $H_i = \sum_{j=1}^i 1/j$ is the i^{th} harmonic number

$$\begin{aligned}1/\alpha * (H_m - H_{m-n}) &= 1/\alpha * \sum_{k=m-n+1}^m 1/k \\ &\leq 1/\alpha * \int_{m-n}^m (1/x) dx \\ &= 1/\alpha * \ln m/(m-n) \\ &= 1/\alpha * \ln 1/(1-\alpha)\end{aligned}$$

table 50% full, 1.387 次

table 90% full, 2.559 次

9 Universal hashing and perfect hashing

全域哈希和完全哈希

Perfect hashing: search time $\Theta(1)$ worst case when the set of keys are static(that is, when the set of keys never change once stored).

9.1 Universal hashing

weakness of hashing

For any choice of hash function, there exists a bad set of keys that all hash to the same slot.

The idea is to choose a hash function at random independent from keys.

This method is called **universal hashing**.

Def: Let U be a universe of keys, and H a finite collection of hash functions mapping U to $0, 1, \dots, m-1$.

Universal collection of hash functions

Such a collection is said to be **universal** if for each pair of distinct keys $k, l \in U$, the number of hash functions $h \in H$ for which $h(k) = h(l)$ is at most $|H|/m$.

In other words, with a hash function randomly chosen from H , $Pr\{h(k) = h(l)\} \leq 1/m$

Thm: if we choose h randomly from the set of hash function H , and then we suppose we are hashing n keys into m slots in table T , then, for a given key x ,

E

Prof:

Let C_x be random variable denoting total number of collisions of keys in T with x , and

$|H|/m$ and $1/m$ 都不理解

let indicator random variable $C_{xy} = 1$ if $h(x) = h(y)$ and 0 otherwise

Note: $E[C_{xy}] = 1/m$ and $C_x = \sum_{y \in T-x} C_{xy}$

After the calculation, we get $E[C_x] = (n-1)/m$

这个概率暂时不懂

Constructing a universal hash function

Let m be prime, decompose key k into $r+1$ digits:

$k = \langle k_0, k_1, \dots, k_r \rangle$ where $0 \leq k_i \leq m-1$

Pick $a = \langle a_0, a_1, \dots, a_r \rangle$ each a_i is chosen randomly from $0, 1, \dots, m-1$

Define $h_a(k) = (\sum_{i=0}^r a_i * k_i) \bmod m$

How big is

Thm: H is universal

Proof: let

Number theory fact 质数的倒数

Let m be prime. For any $z \in Z_m$ (integers mod m) such that $z \neq 0, \exists$ unique $z' \in Z_m$ such that $z * z' = 1 \bmod m$

Example: $m = 7$

z 1 2 3 4 5 6

z' 1 4 5 2 3 6

9.2 Perfect hashing 完全哈希

Hashing can be used to obtain excellent worst-case performance when the set of keys is **static**: **once the keys are stored in the table, the set of keys never changes.**

Some applications naturally have static sets of keys: consider the set of reserved words 保留关键词 in a programming language, or the set of file names on a CD-ROM.

Perfect hashing: search takes $O(1)$ time in the worst case.

Idea: **2 level hashing scheme with universal hashing at each level.** No collisions at level 2.

The first level is essentially the same as for hashing with chaining.

Instead of making a list of the keys hashing to slot j , however, we use a small secondary hash table S_j with an associated hash function h_j . By choosing the hash functions h_j carefully, we can guarantee that there are no collisions at level 2.

9.2.1 Collision-free

Thm: Hash n keys into $m = n^2$ slots using random h in universal H

then $Pr\{\text{number of collisions}\} < 1/2$

Proof: There are C_n^2 pairs of keys that may collide; each pair collide with probability $1/m$ if h is chosen at random from a universal family of H of hash functions.

Let X be a random variable that counts the number of collisions.

When $m = n^2$, the expected number of collisions is

为什么要用概率, 用期望值我觉得是一个更好的选择

$$E[X] = C_n^2/m = (n^2 - n)/(2n^2) < 1/2$$

Applying Markov's inequality, $Pr[X \geq 1] \leq E[X]/1 < 1/2$

Markov's inequality

Proof

For random variable $x \geq 0$, $Pr(x \geq t) \leq E[x]/t$

Proof : $E[x] =$

在上面定理所描述的情况下 $m = n^2$, 随机选择的一个 hash function 很有可能没有 collisions.

所以当 the set K of n keys to be hashed (remember that K is static), it is thus easy to find a collision-free hash function h with a few random trials.

当 n 很大时, hash table $m = n^2$ 占用的空间就会很大, 所以我们采用 2 level hashing approach.

First-level, hash function h hashes the keys into $m = n$ slots

Second-level, if n_j keys hash to slot j , a secondary hash table S_j of size $m_j = n_j^2$ is used to provide collision-free constant-time lookup.

Corollary: $Pr\{\text{no collisions}\} \geq 1/2$

Proof: $Pr \geq 1 \text{ collision} \leq E[\text{number of collisions}]/1 < 1/2$

To find a good level 2 hash function, just test at random. Find one quickly, since $\geq 1/2$ will work.

9.2.2 Analysis of storage

Thm: if we store n keys in hash table of size $m = n$ using a hash function h randomly chosen from a universal class of hash functions, then

$$E\left[\sum_{j=0}^{m-1} g_j^2\right] < 2n$$

where n_j is the number of keys hashing to slot j .

Proof: $a^2 = a + C_a^2$ 恒等式 C_a^2 是一个组合数

We have:

$$\begin{aligned} E\left[\sum_{j=0}^{m-1} n_j^2\right] &= E\left[\sum_{j=0}^{m-1} (n_j + 2 * C_{n_j}^2)\right] \\ &= E\left[\sum_{j=0}^{m-1} n_j\right] + 2 * E\left[\sum_{j=0}^{m-1} C_{n_j}^2\right] \\ &= E[n] + 2 * E\left[\sum_{j=0}^{m-1} C_{n_j}^2\right] \\ &= n + 2 * E\left[\sum_{j=0}^{m-1} C_{n_j}^2\right] \end{aligned}$$

to evaluate the summation $\sum_{j=0}^{m-1} C_{n_j}^2$ we observe it is just the total number of collisions. By the properties of universal hashing, the expected value of this summation is at most

$$C_n^2 * 1/m = n(n-1)/(2m) = (n-1)/2$$

since $m = n$. Thus,

$$E\left[\sum_{j=0}^{m-1} n_j^2\right] \leq n + 2 * (n-1)/2 = 2n - 1 < 2n$$

10 Binary search tree(BST)

Search trees are data structures that support many **dynamic-set operations**, including **search**, **minimum**, **maximum**, **predecessor** 前任, **successor** 继承者, **insert**, **delete**.

Thus, a search tree can be used both as a dictionary and as a priority queue.

These basic operations on a binary search tree take time proportional to the height of the tree.

所以, 对于 complete binary tree with n nodes, these operations run in $\Theta(\lg n)$ worst-case.

10.1 what is a binary search tree?

Each node is an object,

key field and satellite data, and left, right, p(parent) pointers.

if a child or parent is missing, the appropriate field contains the value NIL.

binary search tree property

Let x be a node in a binary search tree.

If y is a node in the left subtree of x , then $key[y] \leq key[x]$.

If y is a node in the right subtree of x , then $key[x] \leq key[y]$.

二叉排序树中, 各结点关键字是惟一的.

注意: 实际应用中, 不能保证被查找的数据集中各元素的关键字互不相同, 所以可将二叉排序树定义中 BST 性质 (1) 里的”小于”改为”小于等于”, 或将 BST 性质 (2) 里的”大于”改为”大于等于”, 甚至可同时修改这两个性质.

10.1.1 tree walk 遍历

Inorder tree walk 中序遍历

the key of the root of a subtree is printed between the values in its left subtree and those in its right subtree.

preorder tree walk: 前序遍历

postorder tree walk 后序遍历

Inorder-tree-walk(x):

```
if  $x \neq \text{nil}$ :
    Inorder-tree-walk(left[ $x$ ])
    print key[ $x$ ]
    Inorder-tree-walk(right[ $x$ ])
```

it takes $\Theta(n)$ time to walk an n -node binary search tree.

10.2 Querying a binary search tree

search, minimum, maximum, successor, predecessor

all these operations $\Theta(h)$ time, where h is the height of the binary search tree

10.2.1 Searching

递归式的

```
Tree-Search(x,k):
    if x=nil or k=key[x]
        return x
    if k < key[x]
        return Tree-Search(left[x], k)
    else return Tree-Search(right[x], k)
```

执行的时候, x 是 root, 表示从 root 开始 search
迭代式的

```
Iterative-Tree-Search(x, k):
    while x != nil and k != key[x]
        if k < key[x]
            x=left[x]
        else x=right[x]
    return x
```

10.2.2 Minimum and maximum

一直往 left subtree 走就可以找到 minimum
一直往 right subtree 走就可以找到 maximum

```
Tree-Minimum(x):
while left[x] != nil
    x=left[x]
return x
```

```
Tree-Maximum(x):
while right[x] != nil
    x=right[x]
return x
```

10.2.3 Successor and predecessor

successor 后继, predecessor 前趋

Successor of a node x is the node with **the smallest key greater than** $key[x]$

A node's in-order successor is the left-most child of its right subtree

后继: 如果没有右子节点, 那么 x 的直接后继结点就是从 x 向上的路径中第一次右转时的节点

Predecessor of a node x is the node with **the greatest key smaller than** $key[x]$

前趋: 如果没有左子节点, 那么 x 的直接前趋节点就是从 x 向上的路径中第一次左转时的节点

A node's in-order predecessor is the right-most child of its left subtree

某个节点有两个子女, 则其后继没有左子女, 其前趋没有右子女.
 因为如果后继有左子女, 由于以后继为 root 的子树在那么左子女

The structure of a binary search tree allows us to determine the successor of a node **without ever comparing keys**.

```

Tree-Successor(x):
if right[x] != nil
    return Tree-Minimum(right[x])
#当右子树非空时, the successor is just the leftmost node in the right subtree

y=p[x] # p is for parent
while y!=nil and x=right[y]
// 这里也可以分成两种情况, x 是p[x] 的left child 与 x 是p[x]的right child 两种情况
    x=y
    y=p[y]
return y
  
```




Figure 15: Binary search tree example

On the other hand, as Exercise 12.2-6 asks you to show,
 if the right subtree of node is empty and x has a successor y , then y is the lowest ancestor of x whose left child is also an ancestor of x . In Figure 15, the successor of the node with key 13 is the node with key 15. To find y , we simply go up the tree from x until we encounter a node that is the left child of its parent; this is accomplished by lines 3-7 of TREE-SUCCESSOR.

Tree-Predecessor is symmetric to Tree-Successor, both run in time $O(h)$

10.3 Insertion and deletion

Insertion

TREE-INSERT(T, z)

```

1  y ← NIL // 这个初始化时必须的, 因为当x一开始就是nil时, 就直接在line 9 对y==nil 进行判断
  
```

```

2  x ← root[T]
3  while x neq NIL
4      do y ← x
5          if key[z] < key[x]
6              then x ← left[x]
7              else x ← right[x]
8  p[z] ← y
9  if y = NIL
10     then root[T] ← z # Tree T was empty
11     else if key[z] < key[y]
12         then left[y] ← z
13         else right[y] ← z

```

line 4: 将 x 保存下来, 是因为 while 循环结束的时候, $x = nil$, 也就是现在 x 是叶子节点的子节点, 这个时候, x 的 parent 就是一个叶子节点, 也就是说 y 是一个叶子节点, 而 z 将要成为 y 的一个子节点.

如果我们不把 x 的 parent 保存下来, 一旦我们走到了一个 leaf 处, 我们就没有办法把 z 放置到我们已经找到的位置.

Deletion

3 cases:

1, if z has **no children**, 见图16, 直接删除掉 z 就行了




Figure 16: Case 1 of deletion

2, if z has only a **single child**, 见图17, 那么删除 z , 将 z 的父节点与此孩子节点 (子树) 关联就可以了




Figure 17: Case 2 of deletion

3, if z has **two children**, 见图18, 删除掉 z 的后继, 然后用 z 的后继来替代 z 因为只有将 z 的后继放在 z 的位置, 才能够保持 BST 的性质, z 的后继由于是在 z 的右子树中的节点, 所以是大于 z 的左子树中所有的节点; 同时后继是右子树中最小的节点, 所以将这个后继挪到 z

的位置后, 是小于右子树中的所有节点




Figure 18: Case 3 of deletion

我们可以将 case 3 化为 case 1 和 case 2 两种已经解决的情况, 我们可以任意选取一个 z 的子树, 从底部开始, 把这个子树上所有的节点, 一个一个从树上摘下来, 存着, 然后 z 就成为一个只有一颗子树的节点, 然后就适用 case 2. 删除完 z 后, 再把存着的那些节点重新插入到树上, 就完成了. 这个方法虽然有点笨, 但是确实是可行的.

其他办法: 要想不影响树的特性, 那么最简单的想法是如果能只动 z 以及 z 以下的节点, 那么树的其它部分肯定是 OK 的. 那只要处理 z 以及 z 的子树就行了. 想想, 如果能在子树中找一个节点来替代 z 的位置, 并保证新的子树也是满足二叉查找树要求的, 这样改动量可能就比较小了. 那么找哪个节点来代替它呢? 当然是键值最接近 X 的, 这样二叉树的特征就比较容易保持嘛. 键值最接近 z 的, 上面已经说过了, 就是直接前趋和直接后继. 正好, 对于有两个子树的 z 来说, 它的直接前趋和直接后继都是在它的子树中的, 分别是左子树的最大值, 右子树的最小值. 而且, 从子树中取下这两个节点 (取下来干嘛? 代替需要删除的 X 节点呗), 也是比较容易的, 因为“最大”“最小”值节点, 最多拥有不超过一个子节点 (不然它绝对不够格做最大或最小). 而没有子节点和只有一个子节点的节点删除, 是我们已经会啦. 好, 那么就取前趋或后继就来代替需要删除的节点, 问题就解决了.

```

TREE-DELETE( $T, z$ )
1  if left[ $z$ ] = NIL or right[ $z$ ] = NIL
2      then  $y \leftarrow z$ 
3      else  $y \leftarrow \text{TREE-SUCCESSOR}(z)$  //也可以用predecessor
4  if left[ $y$ ] neq NIL
5      then  $x \leftarrow \text{left}[y]$ 
6      else  $x \leftarrow \text{right}[y]$ 
7  if  $x$  neq NIL
8      then  $p[x] \leftarrow p[y]$ 
9  if  $p[y] = \text{NIL}$ 
10     then root[ $T$ ]  $\leftarrow x$ 
11     else if  $y = \text{left}[p[y]]$ 
12         then left[ $p[y]$ ]  $\leftarrow x$ 
13         else right[ $p[y]$ ]  $\leftarrow x$ 
14 if  $y$  neq  $z$ 
15     then key[ $z$ ]  $\leftarrow \text{key}[y]$ 
16     copy  $y$ 's satellite data into  $z$ 
17 return  $y$ 

```

这个代码不知道为什么不分三种情况来写, 非要写在一起, 看的就头疼, 还没有看懂

10.4 Randomly built binary search trees

exponential height

11 Red-Black Tress

balanced search tree

$O(\lg n)$ height

Examples:

AVL trees

2-3 trees

2-3-4 trees

B-trees

Red-Black trees

Skip lists

Treaps

11.1 Properties of red-black trees

Red-black trees are one of many search-tree schemes that are "balanced" in order to guarantee that basic dynamic-set operations take $O(\lg n)$ time in the worst case.

red-black properties:

1. Every node is either red or black.
2. The root is black.
3. Every leaf (*NIL*) is black.
4. If a node is red, then both its children are black.
5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes.

4 说明 red 的 parent 必须是 black, 也就是说不能出现连续的两个 red, 同时也说明 at least half the nodes on any simple path from the root to a leaf, not including the root, must be black.

$bh(x)$: black height of x , 不包括 x 节点, 但包括叶节点 $T.nil$

we use the one sentinel $T : nil$ to represent all the NILs—all leaves and the root's parent for saving space.

在实际适用中, 为了方便, 我们有时不将 *nil* 画出.

Lemma: A red-black tree with n internal nodes has height at most $2 \lg(n + 1)$.

Proof: We start by showing that the subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes. 这个用数学归纳法证明

To complete the proof of the lemma, let h be the height of the tree.

the black-height of the root must be at least $h/2$; thus,




Figure 19: Presentation of red black tree

$$n \geq 2^{h/2} - 1 \Rightarrow h \leq 2 \lg(n + 1)$$

Although the algorithms TREE-INSERT and TREE-DELETE from Chapter 12 run in $O(\lg n)$ time

When given a red-black tree as input, they do not directly support the dynamic-set operations INSERT and DELETE, since they do not guarantee that the modified binary search tree will be a red-black tree. 后面会讲到

11.2 Rotations

A local operation in a search tree that preserves the binary-search-tree property tree rotation




Figure 20: Rotation in a search tree

对, 中序遍历没有影响

LEFT-ROTATE(T, x)

```

y=x.right //set y
x.right=y.left //turn y's left subtree into x's right subtree
if y.left != T.nil
    y.left.p = x
y.p = x.p
if x.p == T.nil //link x's parent to y
    T.root = y

```

```

elseif x == x.p.left
    x.p.left = y
else x.p.right = y
y.left = x //put x on y's left
x.p = y

```




Figure 21: Rotation example of a red black tree

11.3 Insertion

当我们 insert 的时候, 我们总是把新的 node color red, 因为如果 color black, 那么 distance(black height) 会成为棘手的问题, 相对于 red-red clash 而言

当进入到 case 3, 表示就结束了

当进入到 case 2 时, 说明很快就会进入 case 3, 所以也表示结束

所以在计算时间的时候, 我们只需要考虑 case 1

在 case 1 的情形里, 每次都会将 violation 往上移动 two generations, 而树的高度 $h \leq 2\lg(n+1)$, 所以情形一的循环次数最多为 $\lg(n+1)$

所以最终是 $\lg n$ time

RB-Insert(T, x):

Tree-Insert(T, x)

color[x] = red

while $x \neq \text{root}[T]$ or color[x] = red

//当 x 是根节点或者 x 是black node(不会有violation), 循环结束

if $p[x] = \text{left}[p[p[x]]]$ //A 类, A 与B是对称的, B: $p[x] = \text{right}[p[p[x]]]$

$y = \text{right}[p[p[x]]]$ //y is x's uncle


```

    if color[y]=red //y is red
        then <case 1>
    else if x=right[p[x]] //bad side
        then <case 2>
        <case 3> //good side,这里不用else, 因为case 2也会进入case 3
    else (same as catogory A with "right" and "left" exchanged) //B
    color[root[T]]=black //改变root的color 对distance 没有影响

```

11.4 Deletion

RB-DELETE(T, z)

```

1 if left[z] = nil[T] or right[z] = nil[T]
2   then y ← z
3   else y ← TREE-SUCCESSOR(z) // y is the realDelNode
4 if left[y] neq nil[T] //if y has no children, x = nil
5   then x ← left[y]
6   else x ← right[y]
7 p[x] ← p[y]
8 if p[y] = nil[T]
9   then root[T] ← x
10  else if y = left[p[y]]
11      then left[p[y]] ← x
12      else right[p[y]] ← x
13 if y 3neq z
14   then key[z] ← key[y]
15       copy y's satelllite data into z
16 if color[y] = BLACK //if red, no need to fix up
17   then RB-DELETE-FIXUP(T, x)
18 return y

```

RB-DELETE-FIXUP(T, x)

```

1 while x neq root[T] and color[x] = BLACK
2   do if x = left[p[x]]
3       then w ← right[p[x]]
4           if color[w] = RED
5               then color[w] ← BLACK                \\ Case 1
6                   color[p[x]] ← RED                \\ Case 1
7                   LEFT-ROTATE(T, p[x])              \\ Case 1
8                   w ← right[p[x]]                    \\ Case 1
9           if color[left[w]] = BLACK and color[right[w]] = BLACK
10              then color[w] ← RED                    \\ Case 2
11              x =p[x]                                \\ Case 2

```

```

12         else if color[right[w]] = BLACK
13             then color[left[w]] ← BLACK           \\ Case 3
14                 color[w] ← RED                     \\ Case 3
15                 RIGHT-ROTATE(T, w)                 \\ Case 3
16                 w ← right[p[x]]                   \\ Case 3
17             color[w] ← color[p[x]]                 \\ Case 4
18             color[p[x]] ← BLACK                     \\ Case 4
19             color[right[w]] ← BLACK                 \\ Case 4
20             LEFT-ROTATE(T, p[x])                   \\ Case 4
21             x ← root[T]                             \\ Case 4
22     else (same as then clause with "right" and "left" exchanged)
23 color[x] ← BLACK

```

Within the while loop, x always points to a nonroot doubly black node

Case 1: x 's sibling w is red See Figure 22

Case 1 (lines 5–8 of RB-DELETE-FIXUP and Figure 22 occurs when node w , the sibling of node




Figure 22: Case 1 for the deletion in a red black tree

x , is red.

Since w must have black children, we can switch the colors of w and $x : p$ and then perform a left-rotation on $x : p$ without violating any of the red-black properties. The new sibling of x , which is one of w 's children prior to the rotation, is now black, and thus we have converted case 1 into case 2, 3, or 4.

Cases 2, 3, and 4 occur when node w is black; they are distinguished by the colors of w 's children.

Case 2: x 's sibling w is black, and both of w 's children are black See Figure 23

In case 2 (lines 10–11 of RB-DELETE-FIXUP and Figure 23), both of w 's children are black. Since




Figure 23: Case 2 for the deletion in a red black tree

w is also black, we take one black off both x and w , leaving x with only one black and leaving w red. To compensate for removing one black from x and w , we would like to add an extra black to $x : p$, which was originally either red or black. We do so by repeating the while loop with $x : p$ as the new node x . Observe that if we enter case 2 through case 1, the new node x is red-and-black, since the original $x : p$ was red. Hence, the value c of the color attribute of the new node x is RED, and the loop terminates when it tests the loop condition. We then color the new node x (singly) black in line 23.

Case 3: x 's sibling w is black, w 's left child is red, and w 's right child is black See Figure 24

Case 3 (lines 13–16 and Figure 24) occurs when w is black, its left child is red, and its right child




Figure 24: Case 3 for the deletion in a red black tree

is black.

We can switch the colors of w and its left child $w:left$ and then perform a right rotation on w without violating any of the red-black properties. The new sibling w of x is now a black node with a red right child, and thus we have transformed case 3 into case 4.

Case 4: x 's sibling w is black, and w 's right child is red See Figure 25

Case 4 (lines 17–21 and Figure 25) occurs when node x 's sibling w is black and w 's right child is




Figure 25: Case 4 for the deletion in a red black tree

red. By making some color changes and performing a left rotation on $x : p$, we can remove the extra black on x , making it singly black, without violating any of the red-black properties. Setting x to be the root causes the while loop to terminate when it tests the loop condition.

11.5 Analysis

height of a red-black tree of n nodes is $O(\lg n)$

the total cost of the procedure without the call to RB-DELETE-FIXUP takes $O(\lg n)$ time.

Within RB-DELETE-FIXUP, each of **cases 1, 3, and 4** lead to **termination** after performing a constant number of color changes and at most three rotations.

Case 2 is the only case in which the while loop can be repeated, and then the pointer x moves up the tree at most $O(\lg n)$ times, performing no rotations. Thus, the procedure RB-DELETE-FIXUP takes $O(\lg n)$ time and performs at most three rotations, and the overall time for RB-DELETE is therefore also $O(\lg n)$

12 Augmenting data structures

12.1 Dynamic order statistics

$select(x, i)$ return i^{th} smallest element in the subtree rooted at x

$rank(x)$ return the rank of a given element in the total ordering of the set

$size[x]$ which is the number of nodes in the subtree rooted at x .




Figure 26: An order statistics tree

An **order-statistic tree** T is simply a red-black tree with additional information (size) stored in each node.

$$size[x] = size[left[x]] + size[right[x]] + 1$$

$$size[nil] = 0$$

在 order-statistic tree 中, 可以有相同的 keys, 但是这样我们就需要重新定义一下 rank

rank: the position at which it would be printed in an inorder walk of the tree.

OS-select(x, i)

$r = size[left[x]] + 1$

 if $i = r$

 return x

 else if $i < r$

 return OS-select(left[x], i)

 else return OS-select(right[x], $i - r$)

$O(\lg n)$ time

OS-RANK(T, x)

```

    r = size[left[x]] + 1
    y = x
    while y != root[T]
        if y = right[p[y]]
            r = r + size[left[p[y]]] + 1
        y = p[y]
    return r

```

$O(\lg n)$ time

maintaining subtree sizes

刚开始的时候, 为一个干净的红黑树添加 size, 需要 $O(\lg n)$ time

在 insert 和 delete 操作的时候, update size 是 $O(1)$ 的操作, 因此 insert 和 delete 仍然是 $O(\lg n)$ 的操作

12.2 How to augment a data structure

Theorem 14.1 (Augmenting a red-black tree)

Let f be an attribute that augments a red-black tree T of n nodes, and suppose that the value of f for each node x depends on only the information in nodes x , $x : left$, and $x : right$, possibly including $x : left : f$ and $x : right : f$. Then, we can maintain the values of f in all nodes of T during insertion and deletion without asymptotically affecting the $O(\lg n)$ performance of these operations.

12.3 Interval tree

这里我们研究的都是闭区间

Given a query interval, we can then quickly find an interval in the set that overlaps it

interval trichotomy

Any two intervals i and i' satisfy the interval trichotomy; that is, exactly one of the following three properties holds:

1. i and i' overlap,
2. i is to the left of i' (i.e., $high[i] < low[i']$),
3. i is to the right of i' (i.e., $high[i'] < low[i]$).

we say that intervals a and b **overlap** if $a \cap b \neq \emptyset$ (Tex 中空集的表示方法), that is if $low[a] \leq high[b]$ and $low[b] \leq high[a]$

An interval tree is a red-black tree that maintains a dynamic set of elements, with each element x containing an interval $x : int$.

INTERVAL-INSERT(T, x)

INTERVAL-DELETE(T, x)

INTERVAL-SEARCH(T, i): returns a pointer to an element x in the interval tree T such that $int[x]$ overlaps interval i , or a pointer to the sentinel $nil[T]$ if no such element is in the set.




Figure 27: An interval tree

Underlying data structure

We choose a red-black tree in which each node x contains an interval $x : int$ and the key of x is the low endpoint

Additional information

In addition to the intervals themselves, each node x contains a value $x : max$, which is the maximum value of any interval endpoint stored in the subtree rooted at x

Maintaining the information

We must verify that insertion and deletion take $O(\lg n)$ time on an interval tree of n nodes. We can determine $x : max$ given interval $x : int$ and the max values of node x 's children:

$$max[x] = max(high[int[x]], max[left[x]], max[right[x]])$$

Developing new operations

INTERVAL-SEARCH(T, i)

```

x = root[T]
while x != nil[T] and i does not overlap int[x]
    if left[x] != nil[T] and max[left[x]] >= low[i]
        x = left[x]
    else x = right[x]
return x

```

也可以写成递归的形式, tail recursive call

list all overlapping intervals:

if I have k overlaps, then the time is $O(k \lg n)$

if I search the second time, I get the same interval, so when I find one, delete it, and search again
output sensitive, 运行时间取决于输出结果

Theorem: let $L = i \in left[x], R = i \in right[x]$

1. if search goes right, then $i' \in L, i'$ overlaps $i = \emptyset$

2. if search goes left, then $i' \in L, i'$ overlaps $i = \emptyset \Rightarrow i' \in R, i'$ overlaps $i = \emptyset$

对于 2. search goes left, 如果这时在左边没有找到 overlaps, 那么我们也不可能在右边找到

segment tree

13 Skip list

dynamic search structure

efficient, randomized, simple

$O(\lg n)$ in expectation with high probability ($\simeq 1 - 1/(n^\alpha)$)

(sorted linked list)

2 sorted linked lists

L2: stores all elements

L1: stores some elements

14 Dynamic programming

optimal sub structure

a very powerful algorithmic paradigm in which a problem is solved by **identifying a collection of subproblems and tackling them one by one, smallest first, using the answers to small problems to help figure out larger ones, until the whole lot of them is solved.**

In dynamic programming we are not given a dag; the dag is implicit. Its nodes are the subproblems we define, and its edges are the dependencies between the subproblems: if to solve subproblem B we need the answer to subproblem A , then there is a (conceptual) edge from A to B . In this case, A is thought of as a smaller subproblem than B — and it will always be smaller, in an obvious sense.

Dynamic programming is effective when a given sub problem may arise from more than one partial set of choices; **the key technique is to store the solution to each such sub problem in case it should reappear.**

A dynamic programming algorithm will examine all possible ways to solve the problem and will pick the best solution. Therefore, we can roughly think of dynamic programming as an intelligent, brute-force method that enables us to go through all possible solutions to pick the best one. If the scope of the problem is such that going through all possible solutions is possible and fast enough, dynamic programming guarantees finding the optimal solution. The alternatives are many, such as using a greedy algorithm, which picks the best possible choice "at any possible branch in the road". While a greedy algorithm does not guarantee the optimal solution, it is faster. Fortunately, some greedy algorithms (such as minimum spanning trees) are proven to lead to the optimal solution.

For example, let's say that you have to get from point A to point B as fast as possible, in a given city, during rush hour.

A dynamic programming algorithm will look into the entire traffic report, looking into all possible combinations of roads you might take, and will only then tell you which way is the fastest. Of course, you might have to wait for a while until the algorithm finishes, and only then can you start driving. The path you will take will be the fastest one (assuming that nothing changed in the external environment). On the other hand, a greedy algorithm will start you driving immediately and

will pick the road that looks the fastest at every intersection.

As you can imagine, this strategy might not lead to the fastest arrival time, since you might take some "easy" streets and then find yourself hopelessly stuck in a traffic jam.

Greedy algorithms: make each choice in a locally optimal manner.

Amortized analysis: Instead of bounding the cost of the sequence of operations by bounding the actual cost of each operation separately, an amortized analysis provides a bound on the actual cost of the entire sequence. One advantage of this approach is that although some operations might be expensive, many others might be cheap.

We typically apply dynamic programming to **optimization problems**. Such problems can have many possible solutions. Each solution has a value, and we wish to **find a solution with the optimal (minimum or maximum) value**. We call such a solution an optimal solution to the problem, as opposed to the optimal solution, since there may be several solutions that achieve the optimal value.

When developing a dynamic-programming algorithm, we follow a sequence of four steps:

1. Characterize the structure of an optimal solution.
2. Recursively define the value of an optimal solution.
3. Compute the value of an optimal solution, typically in a bottom-up fashion.
4. Construct an optimal solution from computed information

14.1 Assembly-line scheduling

一个车间有 2 条生产线, each assembly line with n stations, numbered by $j = 1, 2, 3, \dots, n$

$S_{i,j}$ the j^{th} station on assembly line i ($i = 1$ or 2)

j^{th} station on line 1 and j^{th} station on line 2 have the same function, but the time required varies.

assembly time on line i at station j is $a_{i,j}$

transfer time from line i at station j is $t_{i,j}$, $j = 1, 2, \dots, n-1$

enter time: e_i

exit time: x_i

Step 1: the structure of the fastest way through the factory

let us consider the fastest possible way for a chassis to get from the starting point through station $S_{1,j}$

if $j = 1$, there is only one way

if $j = 2, 3, \dots, n-1$, two ways, come from station $S_{1,j-1}$ and then directly to station $S_{1,j}$

or come from station $S_{2,j-1}$ and be transferred to station $S_{1,j}$

Step 2: a recursive solution

let $f_i[j]$ denote the fastest possible way to get a chassis from the starting point through station $S_{i,j}$

我们最终要求的是: 让一个 chassis 最快速的通过这个组装线, we note it as f^*

$$f^* = \min(f_1[n] + x_1, f_2[n] + x_2)$$

$$f_1[1] = e_1 + a_{1,1}$$

$$f_1[j] = \min(f_1[j-1] + a_{1,j}, f_2[j-1] + t_{2,j-1} + a_{1,j}) \text{ si } j \geq 2$$

$$f_2[1] = e_2 + a_{2,1}$$

$$f_2[j] = \min(f_2[j-1] + a_{2,j}, f_1[j-1] + t_{1,j-1} + a_{2,j}) \text{ si } j \geq 2$$

the $f_i[j]$ values give the values of optimal solutions to sub problems. To help us keep track of how to construct an optimal solution, let us define $l_i[j]$ to be the line number i (1 or 2), whose j^{th} station is used through the station $S_{i,j}$. here, $j = 2, 3, \dots, n$

l^* the line whose station n is used

For example: $l^* = 1$, means we use station $S_{1,6}$

now we look at $l_1[6]$, which is 2, means we use station $S_{2,5}$

now we look at $l_2[5]$, which is 2, means we use station $S_{2,4}$

now we look at $l_2[4]$, which is 1, means we use station $S_{1,3}$

now we look at $l_1[3]$, which is 2, means we use station $S_{2,2}$

now we look at $l_2[2]$, which is 1, means we use station $S_{1,1}$

这样我们就能够逆序追踪我们的最优解

Step 3: Computing the fastest times

如果直接计算, 类似于 Fibonacci 数列的直接计算, 需要 exponential time,

但是也同样类似于 Fibonacci 数列, 后面的值只是取决于前面的值, 所以我们如果把前面已经计算出来的值保存下来, 每个 f 的计算都只需要常数的时间

FASTEST-WAY(a,t,e,x,n)

```

f_1[1] = e_1 + a_{1,1}
f_2[1] = e_2 + a_{2,1}
for j = 2 to n:
    if f_1[j-1] + a_{1,j} <= f_2[j-1] + t_{2,j-1} + a_{1,j}
        f_1[j] = f_1[j-1] + a_{1,j}
        l_1[j] = 1
    else
        f_1[j] = f_2[j-1] + t_{2,j-1} + a_{1,j}
        l_1[j] = 2
    if f_2[j-1] + a_{2,j} <= f_1[j-1] + t_{1,j-1} + a_{2,j}
        f_2[j] = f_2[j-1] + a_{2,j}
        l_2[j] = 2
    else
        f_2[j] = f_1[j-1] + t_{1,j-1} + a_{2,j}
        l_2[j] = 1 //end for

if f_1[n] + x_1 <= f_2[n] + x_2
    f^* = f_1[n] + x_1
    l^* = 1

```

```

else
    f^* = f_2[n] + x_2
    l^* = 2

```

$\Omega(n)$ time

Step 4: Constructing the fastest way through the factory

```

PRINT-STATIONS(l,n)
    i = l^*
    print "line ",i," station ",n
    for j = n downto 2
        i = l_i[j]
        print "line ",i," station ",n

```

这个打印是逆序的, 如果要正序输出, 可以采用 recursive, 或者 stack

14.2 Longest common subsequence (LCS)

Ex: longest common subsequence (LCS)

x : A B C B D A B

y : B D C A B A

LCS: BDAB, BCAB, BCBA 都是长度为 4, 没有长度为 5

Analysis

check if a sequence is in $y : O(n)$, y is of length n

2^m subsequences of x (x is of length m)

check all sub sequences of $x : O(n \cdot 2^m)$

exponential time = slow!

Simplification

1. look at the length of $LCS(x,y)$

2. Extend the alg to find LCS itself

Notation: $|s|$ denotes length of seq s

Strategy: consider **prefixes** of x and y

Define $c[i, j] = |LCS(x[1...i], y[1...j])|$

Then, $c[m, n] = |LCS(x, y)|$

找出 $c[i, j]$ 的通项公式

Theorem:

$$c[i, j] = c[i - 1, j - 1] + 1 \text{ if } x[i] = y[j]$$

$$c[i, j] = \max(c[i, j - 1], c[i - 1, j]) \text{ otherwise}$$

Proof:

14.3 Fibonacci

top-down approach

```
var m := map(0 → 0, 1 → 1)
function fib(n)
  if map m does not contain key n
    m[n] := fib(n - 1) + fib(n - 2)
  return m[n]
```

$O(n)$ time but requires $O(n)$ space top-down approach 也可以减少空间的使用, 就是向 bottom-up 一样, 覆盖掉不用的值

bottom-up approach

```
function fib(n)
  if n = 0
    return 0
  var previousFib := 0, currentFib := 1
  else repeat n - 1 times // loop is skipped if n=1
    var newFib := previousFib + currentFib
    previousFib := currentFib
    currentFib := newFib
  return currentFib
```

$O(n)$ time and requires $O(1)$ space

bottom-up 把 previousFib 之前的都丢了, 所以可以减少空间的使用

14.4 Optimal binary search trees

might be compared with **Huffman trees**, which similarly seek to **place frequently used items near the root in order to produce a dense information encoding**; however, Huffman trees only store data elements in leaves and these elements need not be ordered.

14.5 Common subproblems

Finding the right subproblem takes creativity and experimentation. But there are a few standard choices that seem to arise repeatedly in dynamic programming.

i. The input is x_1, x_2, \dots, x_n and a subproblem is x_1, x_2, \dots, x_i .

The number of subproblems is therefore linear.

ii. The input is x_1, \dots, x_n , and y_1, \dots, y_m . A subproblem is x_1, \dots, x_i and y_1, \dots, y_j .

The number of subproblems is $O(mn)$.

iii. The number of subproblems is $O(mn)$.

The number of subproblems is $O(n^2)$.

iv. The input is a rooted tree. A subproblem is a rooted subtree.

If the tree has n nodes, how many subproblems are there? exponential

15 贪婪算法和最小生成树

A greed algorithm always makes the choice that looks best at the moment. That is, it makes a locally optimal choice in the hope that this choice will lead to a globally optimal solution.

Dynamic programming usually solves the sub problems bottom up, a greedy algorithm usually progresses in a top-down fashion, making one greedy choice after another, reducing each given problem instance to a smaller one.

15.1 Greedy versus dynamic programming

The 0-1 knapsack problem:

Insolvable by greedy strategy, solvable by dynamic programming

Unable to fill the knapsack to capacity, and the empty space lowers the effective value per pound.

The fractional knapsack problem:

Solvable by greedy strategy, compute the value per pound v_i/w_i for each item, take as much as possible of the item with the greatest value per pound.

15.2 Tree

Graph: $E = O(V^2) \Rightarrow \lg E = O(\lg V)$ //省略绝对值的写法

digraph: directed graph

connected $\Rightarrow |E| > |V| - 1$: there is a path from any vertex to any other vertex in the graph.

adjacency matrix 邻接矩阵 $n \times n$ matrix

where $V = 1, 2, \dots, n$

$A[i, j] = 1$ si $(i, j) \in E$

$A[i, j] = 0$ si $(i, j) \notin E$




Figure 28: Represent a graph with adjacency matrix

Represent a graph with adjacency matrix, $O(V^2)$ storage \Rightarrow dense representation

a complete graph, all elements are 1

a sparse graph 稀疏矩阵: 极端例子 a linked list

Adjacency list

矩阵中每一行为 1 组成的一个 list




Figure 29: Represent a graph with adjacency list

$|\text{Adj}[j]| = \text{degree}$ if undirected graph

$|\text{Adj}[j]| = \text{out degree}$ if directed graph

Handshaking Lemma (undirected graph)

$$\sum_{v \in V} \text{degree}(v) = 2|E|$$

每加入一条边, 这条边的两个顶点的 degree 都增加 1

For undirected graphs \Rightarrow adjacency list uses $\Theta(V + E)$ storage

15.3 Huffman code

Huffman coding is a form of statistical coding.

Decoding: Once receiver has tree, it scans incoming bit stream, 0 means go left and 1 means go right.

Huffman coding is an entropy encoding algorithm used for lossless data compression.

Huffman's paper: A method for the construction of minimum redundancy codes

Arithmetic coding and LZW coding

Linear time if input probabilities are sorted

Uniform probability distribution and a number of members which is a power of two: Huffman coding is equivalent to simpler binary block encoding, eg ASCII coding.

Shannon-Fano coding

A symbol with zero probability has zero contribution to the entropy.

$$\lim_{p \rightarrow 0} p \cdot \log_2 p = 0$$

Entropy: the theoretical limit length established by Shannon.

A prefix code (sometimes called "prefix free codes") that is the bit string representing some particular symbol is never a prefix of the bit string representing another symbol.

15.4 MST - Minimum spanning tree

Distributed system

Input: connected, undirected graph $G = (V, E)$

$N : E \rightarrow R$ 权值

For simplicity, assume all edge weights are distinct $\Rightarrow N$ is injective

Output: A spanning tree T (connects all vertexes) of minimum weight

$W(T) = \sum_{(u,v) \in T} w(u,v)$ $// (u,v)$ is an edge




Figure 30: 红色是最小生成树的结果

remove (u,v) from T , then T is partitioned into two sub trees T_1 and T_2

Theorem: T_1 is MST for $G_1 = (V_1, E_1)$ the sub graph of G induced by vertexes in T_1

for T_2 the same thing

Proof: cut and paste

$W(T) = W(u,v) + W(T_1) + W(T_2)$

if $\exists T'$ better than T_1 for G_1 , then

overlapping sub problems? Yes

dynamic programming? Yes, but MST exhibits a better property \Rightarrow greedy algorithm

Theorem: Let T be MST of $G = (V, E)$, Suppose (u,v) in E , least-weight edge connecting A to $V - A$, then (u,v) in T

Proof: cut and paste

16 Amortized Analysis

16.1 Incrementing a binary counter

A k -bit binary counter that counts upward from 0.

Use an array $A[0..k-1]$ of bits, where $A.length = k$, as the counter.

A binary number x that is stored in the counter has its lowest-order bit in $A[0]$ and its highest-order bit in $A[k-1]$, so that $x = \sum_{i=0}^{k-1} A[i] * 2^i$. Initially, $x = 0$

Increment(A)

```

i = 0
while i < A.length and A[i] == 1
    A[i] = 0
    i = i + 1
if i < A.length
    A[i] = 1

```

A single execution of INCREMENT takes time $\Theta(k)$ in the worst case, in which array A contains all 1s. Thus, a sequence of n INCREMENT operations on an initially zero counter takes time $O(nk)$ in the worst case.

We can tighten our analysis to yield a worst-case cost of $O(n)$ for a sequence of n INCREMENT operations by observing that not all bits flip each time INCREMENT is called.

$A[0]$ does flip each time INCREMENT is called

$A[1]$, flips only every other time, flips $n/2$ times

Similarly, bit $A[2]$ flips only every fourth time, or $n/4$ times

bit $A[i]$ flips $n/(2^i)$ times in a sequence of n INCREMENT operations on an initially zero counter.

For $i \geq k$, bit $A[i]$ does not exist, and so it cannot flip. The total number of flips in the sequence is thus

$$\sum_{i=0}^{k-1} \lfloor \frac{n}{2^i} \rfloor < n \sum_{i=0}^{\infty} \frac{1}{2^i} = 2n$$

The worst-case time for a sequence of n INCREMENT operations on an initially zero counter is therefore $O(n)$. The average cost of each operation, and therefore the amortized cost per operation, is $O(n)/n = O(1)$

16.2 Dynamic tables

insert elements

```

i 1 2 3 4 5 6 7 8 9 10
size_i 1 2 4 4 8 8 8 8 16 16
cost_i 1 2 3 1 5 1 1 1 9 1

```

Cost of n inserts

$$= \sum_{i=1}^n \text{cost}_i = n + \sum_{j=0}^{\lfloor \lg(n-1) \rfloor} 2^j \leq 3n = \Theta(n)$$

The average cost of a insert = $\Theta(n)/n = \Theta(1)$

Amortized analysis

Analyze a sequence of operations to show that average cost is small, even though 1 or several operations may be expensive

这里不涉及到概率分析

Types of amortized arguments

Aggregate (just saw)

accounting

potential

16.3 Accounting method

charge the i^{th} operation a fictitious amortized cost c_i (1 pays for 1 unit of operation)
 fee is consumed to perform op
 unused amount stored in bank for use in later ops
 bank balance must not go negative: $\sum_{i=1}^n \text{true-cost}_i \leq \sum_{i=1}^n \text{amortized-cost}_i$

Dynamic table

Charge $c_i = 3\$$ for i^{th} insert

1\$ pays for immediate insert

2\$ stored for table doubling

When table doubles

1\$ moves recent item

1\$ moves old item

0 0 0 0 2 2 2 2

第五项插入时, 收取 3\$, 其中 1\$ 用于自己的插入, 剩下的 2\$ 存入 bank, 第 6, 7, 8 项插入时, 进行同样的操作, 这时银行刚好有 8\$ 的存款

插入第 9 项时, table 需要 double, 银行中剩余的 8\$ 刚好可以支付原来 8 个 elements 的移动到新 table 的花费

0 0 0 0 0 0 0 0 2

Invariant bank balance ≥ 0

$$\sum_{i=1}^n \text{amortized-cost}_i = 3n$$

实际上第一项可以只支付 2\$, 从而可以省下 1\$.

16.4 Potential method

Bank account viewed as potential energy of dynamic set

Framework start with data structure D_0 op i transforms D_{i-1} to D_i

Define potential function

$\Phi : \{D\} \rightarrow R$ such that $\Phi(D_0) = 0$ and $\Phi(D) \geq 0$

Amortized cost ac_i with respect Φ is

$$ac_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$$

if $\Phi(D_i) - \Phi(D_{i-1}) > 0$ overcharge

if $\Phi(D_i) - \Phi(D_{i-1}) < 0$ undercharge

$$\sum_{i=0}^n ac_i = \sum_{i=0}^n (c_i + \Phi(D_i) - \Phi(D_{i-1})) = \sum_{i=0}^n c_i + \Phi(D_n) - \Phi(D_0) = \sum_{i=0}^n c_i + \Phi(D_n)$$

我们需要 $\sum_{i=0}^n ac_i \geq \sum_{i=0}^n c_i$, 所以应该保证 $\Phi(D_n) \geq 0$, 但是实际上我们不知道应该怎样选择这样一个 potential function, 所以我们通常保证 $\Phi(D_i) \geq 0$

Table doubling

Define: $\Phi(D_i) = 2i - 2^{\text{ceil}(\lg i)}$

(Assume $2^{\text{ceil}(\lg 0)} = 0$)

Note: $\Phi(D_0) = 0$ and $\Phi(D_i) \geq 0$ (because $\text{ceil}(\lg i) = \lg i$ or $\lg i + 1$)

For example:

full full full full full full

插入了 6 个, 此时 $\Phi = 2 * 6 - 2^{\text{ceil}(\lg 6)} = 12 - 8 = 4$

如果我们用 accounting method, 那么

0 0 0 0 2 2

此时 bank balance 是 4, 和上面的 potential 一样

Amortized cost of i^{th} insert

$$\begin{aligned} ac_i &= c_i + \Phi(D_i) - \Phi(D_{i-1}) \\ &= \begin{cases} i & \text{si } i-1 \text{ exact power of } 2 \\ 1 & \text{si non} \end{cases} + (2i - 2^{\text{ceil}(\lg i)}) - (2(i-1) - 2^{\text{ceil}(\lg(i-1))}) \\ &= \begin{cases} i & \text{si } i-1 \text{ exact power of } 2 \\ 1 & \text{si non} \end{cases} + 2 - 2^{\text{ceil}(\lg i)} + 2^{\text{ceil}(\lg(i-1))} \end{aligned}$$

Case 1: $i-1$ is exact power of 2

$$ac_i = i + 2 - 2^{\text{ceil}(\lg i)} + 2^{\text{ceil}(\lg(i-1))} = i + 2 - 2^{\lg(i-1)+1} + 2^{\lg(i-1)} = i + 2 - 2 * (i-1) + (i-1) = 3$$

Case 2: $i-1$ is not exact power of 2

$$ac_i = 1 + 2 - 2^{\text{ceil}(\lg i)} + 2^{\text{ceil}(\lg(i-1))} = 3$$

$\text{ceil}(\lg i)$ 与 $\text{ceil}(\lg(i-1))$ 是相等的

Conclusions

Amortized costs provide a clean abstraction for data structure performance.

Diff potential functions or accounting costs may yield diff bounds

17 Advanced data structures

Because disks operate much more slowly than random-access memory, we measure the performance of B-trees not only by how much computing time the dynamic-set operations consume but also by how many disk accesses they perform.

Fibonacci heaps are key components of some of the asymptotically fastest algorithms to date for graph problems.

Splay trees

Persistent data structure: a data structure that always preserves the previous version of itself when it is modified. Such data structures are effectively immutable, as their operations do not (visibly) update the structure in-place, but instead always yield a new updated structure. (A persistent data structure is not a data structure committed to persistent storage, such as a disk; this is a different and unrelated sense of the word "persistent.")

A data structure is partially persistent if all versions can be accessed but only the newest version can be modified. The data structure is fully persistent if every version can be both accessed and modified. If there is also a meld or merge operation that can create a new version from two previous versions, the data structure is called confluent persistent. Structures that are not persistent are called ephemeral.

Fusion trees

Exponential search trees

18 B-Trees

B-trees are similar to red-black trees (Chapter 13), but they are better at minimizing disk I/O operations. Many database systems use B-trees, or variants of B-trees, to store information.

B-tree nodes may have many children, from a few to thousands.

If an internal B-tree node x contains $x.n$ keys, then x has $x.n+1$ children. The keys in node x serve as dividing points separating the range of keys handled by x into $x.n + 1$ sub ranges, each handled by one child of x .

In a typical B-tree application, the amount of data handled is so large that all the data do not fit into main memory at once. The B-tree algorithms copy selected pages from disk into main memory as needed and write back onto disk the pages that have changed. B-tree algorithms keep only a constant number of pages in main memory at any time; thus, the size of main memory does not limit the size of B-trees that can be handled.

B-tree algorithm depends primarily on the number of DISK-READ and DISK-WRITE operations, we typically want each to read or write as much information as possible. Thus, a B-tree node is usually as large as a whole disk page, and this size limits the number of children a B-tree node can have.

这一段话不懂, 多读取信息和节点的大小有什么关系, 读取信息是物理上的

18.1 Definition of B-Trees

19 AVL tree

Self-balancing binary search tree

The heights of the two child subtrees of any node differ by at most one;

Time complexity in big O notation

	Average	Worst case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(\log n)$
Insert	$O(\log n)$	$O(\log n)$
Delete	$O(\log n)$	$O(\log n)$

Worst case is when right subtree has height 1 more than the left subtree for every node.

N_h : min # of nodes of height h

$$N_h = 1 + N_{h-1} + N_{h-2}$$

It looks like Fibonacci

$N_h > F_h = \frac{\phi^h}{\sqrt{5}}$ with $\phi = 1.618$ golden ratio

$\frac{\phi^h}{\sqrt{5}} < n \Rightarrow h < c * \lg n$

Insert

1. simple BST insert
2. fix AVL property from changed node up

AVL sort

1. insert n items - $\Theta(n \lg n)$
2. in-order traversal $\Theta(n)$

20 Graph

We adopt a common notational convention: only inside asymptotic notation, such as O -notation, Θ -notation, the symbol V denotes $|V|$ and the symbol E denotes $|E|$.

The vertex set of a graph: $V[G]$

The edge set of a graph: $E[G]$

Sparse graph: $|E|$ is much less than $|V|^2$

Breadth-first search: BFS

Breadth-first search is so named because it expands the frontier between discovered and undiscovered vertices uniformly across the breadth of the frontier. That is, the algorithm discovers all vertices at distance k from s before discovering any vertex at distance $k + 1$.

三种颜色的含义:

White: not having been discovered

Gray: frontier, discovered vertices that have not yet had their adjacency lists fully examined

Black: all its adjacency vertices have been examined

$O(V + E)$ time

Breadth-first trees

The procedure BFS builds a breadth-first tree as it searches the graph, the tree is represented by the π field in each vertex.

$G = (V, E)$ with source s

Predecessor subgraph of G as $G\pi = (V\pi, E\pi)$, where

$V\pi = v \in V : \pi[v] \neq \text{NIL} \cup s$

$E\pi = (\pi[v], v) : v \in V\pi - s$

21 最短路径算法

Negative edge weight \Rightarrow some shortest path may not exist.

Ex:

1. 图中出现了一个环, 而且这个环的总权值是负的 negative cycle, 那么从 u 到 v 的路径到达这个环

时, 由于这个环的权值是负的, 它会永远地绕着这个环转下去, 使路径不断地变短

2. u 和 v 之间根本就没有路径

optimal substructure

a sub-path of a shortest path is shortest path. // 用反证法 contradiction

Triangle inequality

$dist(u, v) \leq dist(u, x) + dist(x, v)$ //最短路径的定义

dist 表示最短路径

Single source shortest path problem

从源点 s 到所有其他点的最短路径

Assume all the edge weight are non-negative.

Idea: greedy

21.1 Dijkstra's algorithm

详细的可以看 en wiki

a graph search algorithm that solves the single-source shortest path problem for a graph with non-negative edgepath costs, producing a shortest path tree. This algorithm is often used in routing and as a subroutine in other graph algorithms.

It picks the unvisited vertex with the lowest-distance, calculates the distance through it to each unvisited neighbor, and updates the neighbor's distance if smaller. Mark visited (set to red) when done with neighbors.

这里有一个图片 (algo_Dijkstra.gif), 但是由于图片是 gif 动态的, 无法放在这里

```

1 function Dijkstra(Graph, source):
2     for each vertex v in Graph:                                // Initializations
3         dist[v] := infinity ;    // Unknown distance function from source to v
5         previous[v] := undefined ;    // Previous node in optimal path
6     end for                                                    // from source
7
8     dist[source] := 0 ;    // Distance from source to source
9     Q := the set of all nodes in Graph ; // All nodes in the graph are unoptimized -
                                         //thus are in Q, priority queue
10    S := empty set    // no node optimized - thus S empty
11    while Q is not empty:    // The main loop
12        u := vertex in Q with smallest distance in dist[] ;    // Source node in first case
13        remove u from Q ;
14        if dist[u] = infinity:
15            break ;    // all remaining vertices are
16        end if    // inaccessible from source
17        put u in S
18        for each neighbor v of u: //where v has not yet been removed from Q.
19            alt := dist[u] + dist_between(u, v) ;
20            if alt < dist[v]:    // Relax (u,v,a)
```

```

22         dist[v] := alt ;
23         previous[v] := u ;
24         decrease-key v in Q;      // Reorder v in the Queue
                                     //因为v的距离被重新赋值了，所以它在Q中的值也要

```

相应地降低

```

25         end if
26     end for
27 end while
28 return dist;
29 end function

```

Read the shortest path from source to target by reverse iteration

```

1  S := empty sequence
2  u := target
3  while previous[u] is defined:      // Construct the shortest path with a stack S
4      insert u at the beginning of S      // Push the vertex into the stack
5      u := previous[u]                  // Traverse from target to source
6  end while ;

```

Shortest path tree: the union of all shortest paths

breath first search: BFS 广度优先搜索

BFS 就是 Dijkstra

two changes

1. 不使用 priority queue, 而是使用 FIFO queue
2. relax

22 String matching

The Rabin-Karp algorithm

String matching with finite automata

The Knuth-Morris-Pratt algorithm

23 并行

parallel algorithms

dynamic multithreading

share memory, multicore

```

Ex: fib(n)
if n<2
    return n
x = spawn fib(n-1)

```

```

y = spawn fib(n-2)
sync
return (x+y)

```

spawn: 衍生: Subroutine can execute at the same time with parent

Sync: wait until all children are done

caching

24 Appendix

24.1 Summations

harmonic series

$$H_n = 1 + 1/2 + 1/3 + \dots + 1/n = \sum 1/k = \ln n + O(1)$$

24.1.1 Approximation by integrals

When a summation can be expressed as $\sum_{k=m}^n f(k)$, where $f(k)$ is a monotonically increasing function, we can approximate it by integrals:

$$\int_{m-1}^n f(x)dx \leq \sum_{k=m}^n f(k) \leq \int_m^{n+1} f(x)dx$$





Figure 31: approximation by lower bound

Figure 32: approximation by upper bound

when f is monotonically decreasing function:

$$\int_m^{n+1} f(x)dx \leq \sum_{k=m}^n f(k) \leq \int_{m-1}^n f(x)dx$$

n^{th} harmonic number: for a lower bound, we obtain

$$\sum_{k=1}^n \frac{1}{k} \geq \int_1^{n+1} \frac{1}{x} dx = \ln(n+1)$$

for the upper bound, we derive the inequality

$$\sum_{k=2}^n \frac{1}{k} \leq \int_1^n \frac{1}{x} dx = \ln n$$

which yields the bound

$$\sum_{k=1}^n \frac{1}{k} \leq \ln n + 1$$

24.2 Relations

A binary relation R on two sets A and B is a subset of the Cartesian product $A \times B$. If $(a, b) \in R$, we sometimes write $a R b$. When we say that R is a binary relation on a set A , we mean that R is a subset of $A \times A$.

For example, the "less than" relation on the natural numbers is the set $\{(a, b) : a, b \text{ in } N \text{ and } a < b\}$.

An n -ary relation on sets A_1, A_2, \dots, A_n is a subset of $A_1 \times A_2 \times \dots \times A_n$.

A binary relation $R \subseteq A \times A$ is

reflexive if $a R b$ implies $b R a$ for all $a, b \in A$

symmetric if $a R b$ implies $b R a$ for all $a, b \in A$. For example, "="

transitive if $a R b$ and $b R c$ imply $a R c$ for all $a, b, c \in A$.

A relation that is reflexive, symmetric, and transitive is an equivalence relation

Théorème 24.1. An equivalence relation is the same as a partition

The equivalence classes of any equivalence relation R on a set A form a partition of A , and any partition of A determines an equivalence relation on A for which the sets in the partition are the equivalence classes.

A binary relation R on a set A is antisymmetric if $a R b$ and $b R a$ imply $a = b$.

A relation that is reflexive, antisymmetric, and transitive is a partial order 偏序.

And we call a set on which a partial order is defined a partially ordered set.

For example, the relation "is a descendant(子孙, 后裔) of" is a partial order on these of all people (if we view individuals as being their own descendants).

In a partially ordered set A , there may be no single "maximum" element a such that $b R a$ for all $b \in A$. Instead, there may several maximal elements a such that for no $b \in A$, where $b \neq a$, is it the case that $a R b$. For example, in a collection of different-sized boxes there may be several maximal boxes that don't fit inside any other box, yet no single "maximum" box into which any other box will fit.

A partial order R on a set A is a total or linear order if for all $a, b \in A$, we have $a R b$ or $b R a$. That is, every pairing of elements of A can be related by R .

1. The subset relation " \subseteq " on all subsets of Z is a partial order but not a total order.

2. Give examples of relations that are

- reflexive and symmetric but not transitive,
- reflexive and transitive but not symmetric, \leq
- symmetric and transitive but not reflexive.

24.2.1 Functions

Given two sets A and B , a function f is a binary relation on $A \times B$ such that for all $a \in A$, there exists precisely one $b \in B$ such that $(a, b) \in f$

$$f : A \rightarrow B,$$

if $b = f(a)$, a is called the argument of f and b the value of f at a .

When the domain of a function f is a Cartesian product, we often omit the extra parentheses surrounding the argument of f .

eg, $f : A_1 \times A_2 \times \cdots \times A_n \rightarrow B$, we write $b = f(a_1, a_2, \dots, a_n)$ instead of $b = f((a_1, a_2, \dots, a_n))$.

We also call each a_i an argument to the function f , though technically the (single) argument to f is the n -tuple (a_1, a_2, \dots, a_n) .

严格上来讲 (a_1, a_2, \dots, a_n) 是一个参数

Give a bijection: $Z \rightarrow Z \times Z$.

24.3 Graphs

Directed and undirected

A directed graph (or digraph) G is a pair (V, E) , where V is a finite set and E is a binary relation on V .

The set V is called the *vertex* set of G , and its elements are called vertices (singular: vertex).

The set E is called the edge set of G , and its elements are called edges.

In a directed graph, self-loops—edges from a vertex to itself are possible.

Vertices are represented by circles in the figure, and edges are represented by arrows

In an undirected graph $G = (V, E)$, the edge set E consists of unordered pairs of vertices, rather than ordered pairs.

In an undirected graph, self-loops are forbidden.

If (u, v) is an edge in a directed graph $G = (V, E)$, we say that (u, v) is incident from u or leaves vertex u and is incident to or enters vertex v .

If (u, v) is an edge in an undirected graph $G = (V, E)$, we say that (u, v) is incident on vertices u and v .

If (u, v) is an edge in a graph $G = (V, E)$, we say that vertex v is adjacent 毗邻 to vertex u .

The degree of a vertex in an undirected graph is the number of edges incident on it.

A vertex whose degree is 0, such as vertex 4 is isolated.

In a directed graph:

The out-degree of a vertex is the number of edges leaving it,

The in-degree of a vertex is the number of edges entering it.

The degree of a vertex in a directed graph is its in-degree plus its out-degree.

A path of length k from a vertex u to a vertex u' in a graph $G = (V, E)$ is a sequence $\langle v_0, v_1, \dots, v_k \rangle$ of vertices such that $u = v_0, u' = v_k$, and $(v_{i-1}, v_i) \in E$ for $i = 1, 2, \dots, k$.

The length of the path is the number of edges in the path.

If there is a path p from u to u' , we say that u' is reachable from u via p .

A path is simple if all vertices in the path are distinct.

In a directed graph, a path $\langle v_0, v_1, \dots, v_k \rangle$ forms a cycle if $v_0 = v_k$ and the path contains at least one edge.

The cycle is simple if, in addition, v_1, v_2, \dots, v_k are distinct.

A self-loop is a cycle of length 1.

Two paths $\langle v_0, v_1, v_2, \dots, v_{k-1}, v_0 \rangle$ and $\langle v'_0, v'_1, v'_2, \dots, v'_{k-1}, v'_0 \rangle$ form the same cycle if there exists an integer j such that $v'_i = v_{(i+j) \bmod k}$ for $i = 0, 1, \dots, k-1$.

例如:

$v_0, v_1, v_2, v_3, v_4, v_0$

$v'_0, v'_1, v'_2, v'_3, v'_4, v'_0$

如果 $v'_0 = v_3$, 那么

$v'_1 = v_4 = v_{1+3}$

$v'_2 = v_{2+3} = v_5$ 但是实际上应该为 v_0 , 所以这里应该取 $(2+3) \bmod 5$

$v'_3 = v_{(3+3) \bmod 5} = v_1$

$v'_4 = v_{(4+3) \bmod 5} = v_2$

A directed graph with no self-loops is simple. In an undirected graph, a path v_0, v_1, \dots, v_k forms a (simple) cycle if $k \geq 3$, $v_0 = v_k$ and v_1, v_2, \dots, v_k are distinct.

A graph with no cycles is acyclic.

An undirected graph is connected if every pair of vertices is connected by a path.

The connected components of a graph are the equivalence classes of vertices under the "is reachable from" relation.

An undirected graph is connected if it has exactly one connected component, that is, if every vertex is reachable from every other vertex.

A directed graph is strongly connected if every two vertices are reachable from each other. The strongly connected components of a directed graph are the equivalence classes of vertices under the "are mutually reachable" relation. A directed graph is strongly connected if it has only one strongly connected component.

Two graphs $G = (V, E)$ and $G' = (V', E')$ are isomorphic if there exists a bijection $f : V \rightarrow V'$ such that $(u, v) \in E$ if and only if $(f(u), f(v)) \in E'$.

In other words, we can relabel the vertices of G to be vertices of G' , maintaining the corresponding edges in G and G' .




Figure 33: graph isomorphic

$V = 1, 2, 3, 4, 5, 6$ and $V' = u, v, w, x, y, z$.

Bijection $f(1) = u, f(2) = v, f(3) = w, f(4) = x, f(5) = y, f(6) = z$

图33中 a 表示了一个 **isomorphic**, b 没有

Given an undirected graph $G = (V, E)$, the **directed version** of G is the directed graph $G' = (V, E')$, where $(u, v) \in E'$ if and only if $(u, v) \in E$.

Given a directed graph $G = (V, E)$, the **undirected version** of G is the undirected graph $G' = (V, E')$, where $(u, v) \in E'$ if and only if $u \neq v$ and $(u, v) \in E$.

In a directed graph $G = (V, E)$, a **neighbor** of a vertex u is any vertex that is adjacent to u in the undirected version of G

That is, v is a neighbor of u if either $(u, v) \in E$ or $(v, u) \in E$.

In an undirected graph, u and v are neighbors if they are adjacent.

Several kinds of graphs are given special names.

A **complete graph** is an undirected graph in which every pair of vertices is adjacent.

A **bipartite 双边 graph** is an undirected graph $G = (V, E)$ in which V can be partitioned into two sets V_1 and V_2 such that $(u, v) \in E$ implies either $u \in V_1$ and $v \in V_2$ or $u \in V_2$ and $v \in V_1$.

That is, all edges go between the two sets V_1 and V_2 .

An acyclic, undirected graph is a **forest**, and a connected, acyclic, undirected graph is a (free) **tree**. We often take the first letters of "directed acyclic graph" 定向非循环图 and call such a graph a **dag**.

EX: Show that any connected, undirected graph $G = (V, E)$ satisfies $|E| \geq |V| - 1$.

24.3.1 Trees

A free tree is a connected, acyclic, undirected graph. We often omit the adjective "free" when we say that a graph is a tree.

If an undirected graph is acyclic but possibly disconnected, it is a forest.

Many algorithms that work for trees also work for forests.




Figure 34: forest and tree

图34中 a 是 tree, b 是 forest, c 都不是

Théorème 24.2 (Properties of free trees). *Let $G = (V, E)$ be an undirected graph. The following statements are equivalent.*

1. G is a free tree.
2. Any two vertices in G are connected by a unique simple path.

3. G is connected, but if any edge is removed from E the resulting graph is disconnected.
4. G is connected, and $|E| = |V| - 1$.
5. G is acyclic, and $|E| = |V| - 1$.
6. G is acyclic, but if any edge is added to E the resulting graph contains a cycle.

2 中, unique 排除了 cycle 的存在, 因为如果有 cycle, cycle 中的两点间的 path 就不是 unique

3 中, edge 移走后, 点还在

6 中, 只增加 edge, 而不增加点

A **rooted tree** is a free tree in which one of the vertices is distinguished from the others. The distinguished vertex is called the **root** of the tree. We often refer to a vertex of a rooted tree as a **node** of the tree.

Consider a node x in a rooted tree T with root r

Any node y on the unique path from r to x is called an **ancestor** of x

If y is an ancestor of x , then x is a **descendant** of y

(Every node is both an ancestor and a descendant of itself.)

If y is an ancestor of x and $x \neq y$, then y is a **proper ancestor** of x and x is a **proper descendant** of y . The subtree rooted at x is the tree induced by descendants of x , rooted at x .

If the last edge on the path from the root r of a tree T to a node x is (y, x) , then y is the **parent** of x , and x is a **child** of y . The root is the only node in T with no parent. If two nodes have the same parent, they are **siblings**. A node with no children is an **external node or leaf**. A non leaf node is an **internal node**.

The number of children of a node x in a rooted tree T is called the **degree** of x .

The length of the path from the root r to a node x is the **depth** of x in T .

The **height** of a node in a tree is the number of edges on the longest simple downward path from the node to a leaf,

and the height of a tree is the height of its root.

An **ordered tree** is a rooted tree in which the children of each node are ordered. That is, if a node has k children, then there is a first child, a second child... and a k^{th} child.

Binary and positional trees

Binary trees are defined recursively. A **binary tree** T is a structure defined on a finite set of nodes that either

contains no nodes, or

is composed of three disjoint sets of nodes: a **root** node, a binary tree called its **left subtree**, and a binary tree called its **right subtree**.

The binary tree that contains no nodes is called the empty tree or null tree, sometimes denoted NIL.

If the left subtree is nonempty, its root is called the **left child** of the root of the entire tree. Likewise, the root of a non null right subtree is the right child of the root of the entire tree.

Full binary tree

each node is either a leaf or has degree exactly 2. There are no degree-1 nodes.

A k -ary tree is a positional tree in which for every node, all children with labels greater than k are missing.

Thus, a binary tree is a k -ary tree with $k = 2$.

A complete k -ary tree is a k -ary tree in which all leaves have the same depth and all internal nodes have degree k .

How many leaves does a complete k -ary tree of height h have?

The root has k children at depth 1, each of which has k children at depth 2, etc.

Thus, the number of leaves at depth h is k^h .

Consequently, the height of a complete k -ary tree with n leaves is $\log_k n$. The number of internal nodes of a complete k -ary tree of height h is

$$1 + k + k^2 + \cdots + k^{h-1} = \sum_{i=0}^{h-1} k^i = \frac{k^h - 1}{k - 1}$$

Thus, a complete binary tree ($k = 2$) has $2^h - 1$ internal nodes.




Figure 35: tree height

[Exercises B.5-3] Show by induction that the number of degree-2 nodes in any nonempty binary tree is 1 less than the number of leaves.

[Exercises B.5-4] Use induction to show that a nonempty binary tree with n nodes has height at least $\lg n$.

24.4 Probability

The conditional probability of an event A given that another event B occurs is defined to be. 在已知 B 发生的情况下, A 发生的概率.

$$Pr(A|B) = \frac{Pr(A \cap B)}{Pr(B)}$$

whenever $Pr\{B\} \neq 0$. (We read " $Pr\{A|B\}$ " as "the probability of A given B ."). Intuitively, since we are given that event B occurs, the event that A also occurs is $A \cap B$. That is, $A \cap B$ is the set

of outcomes in which both A and B occur. Since the outcome is one of the elementary events in B , we normalize the probabilities of all the elementary events in B by dividing them by $Pr\{B\}$, so that they sum to 1.

Two events are **independent** if $Pr(A \cap B) = Pr(A) \cdot Pr(B)$ which is equivalent, if $Pr\{B\} \neq 0$, to the condition $Pr\{A|B\} = Pr\{A\}$. 也就是说 A 的发生与 B 没有关系

Example 1. *suppose that two fair coins are flipped and that the outcomes are independent. Then the probability of two heads is $(1/2)(1/2) = 1/4$. Now suppose that one event is that the first coin comes up heads and the other event is that the coins come up differently. Each of these events occurs with probability $1/2$, and the probability that both events occur is $1/4$; thus, according to the definition of independence, the events are independent-even though one might think that both events depend on the first coin.*

Finally, suppose that the coins are welded together so that they both fall heads or both fall tails and that the two possibilities are equally likely. Then the probability that each coin comes up heads is $1/2$, but the probability that they both come up heads is $1/2 \neq (1/2)(1/2)$. Consequently, the event that one comes up heads and the event that the other comes up heads are not independent.

A collection A_1, A_2, \dots, A_n of events is said to be **pairwise independent** if

$$Pr\{A_i \cap A_j\} = Pr\{A_i\}Pr\{A_j\}$$

for all $1 \leq i < j \leq n$.

We say that the events of the collection are **(mutually) independent** if every k -subset $A_{i_1}, A_{i_2}, \dots, A_{i_k}$ of the collection, where $2 \leq k \leq n$ and $1 \leq i_1 < i_2 < \dots < i_k \leq n$, satisfies

$$Pr\{A_{i_1} \cap A_{i_2} \cap \dots \cap A_{i_k}\} = Pr\{A_{i_1}\}Pr\{A_{i_2}\} \dots Pr\{A_{i_k}\}$$

24.4.1 Baye's theorem

$$Pr(A \cap B) = \frac{Pr(A) \times Pr(B \cap A)}{Pr(B)}$$

$$B = (B \cap A) \cup (B \cap \bar{A})$$

$B \cap A$ and $B \cap \bar{A}$ are mutually exclusive events

$$Pr(B) = Pr(B \cap A) + Pr(B \cap \bar{A}) = Pr(A) \times Pr(B|A) + Pr(\bar{A}) \times Pr(B|\bar{A})$$

上面的式子可以化为:

$$Pr(A \cap B) = \frac{Pr(A) \times Pr(B \cap A)}{Pr(A) \times Pr(B|A) + Pr(\bar{A}) \times Pr(B|\bar{A})}$$

24.4.2 Discrete random variables

We define two random variables X and Y to be independent if for all x and y , the events $X = x$ and $Y = y$ are independent or equivalently, if for all x and y , we have $Pr(X = x \text{ and } Y = y) = Pr(X = x) \times Pr(Y = y)$.

Linearity of expectation $E[X + Y] = E[X] + E[Y]$

whenever $E[X]$ and $E[Y]$ are defined. And it holds even if X and Y are not independent.

Linearity of expectation is the key property that enables us to perform probabilistic analysis by using indicator random variables.

When two random variables X and Y are independent and each has a defined expectation,

$$\begin{aligned}
 E[XY] &= \sum_x \sum_y x \times y \times Pr(X = x \text{ and } Y = y) \\
 &= \sum_x \sum_y x \times y \times Pr(X = x) \times Pr(Y = y) \\
 &= \left(\sum_x x \times Pr(X = x) \right) \left(\sum_y y \times Pr(Y = y) \right) \\
 &= E[X]E[Y]
 \end{aligned} \tag{24-4-1}$$

In general, when n random variables X_1, X_2, \dots, X_n are mutually independent,

$$E[X_1 \times X_2 \times \dots \times X_n] = E[X_1] \times E[X_2] \times \dots \times E[X_n]$$

When a random variable X takes on values from the set of natural numbers $N=0,1,2, \dots$, there is a nice formula for its expectation:

$$\begin{aligned}
 E[X] &= \sum_{i=0}^{\infty} i \times Pr(X = i) \\
 &= \sum_{i=0}^{\infty} i \times (Pr(X \geq i) - Pr(X \geq i+1)) \\
 &= \sum_{i=1}^{\infty} Pr(X \geq i)
 \end{aligned} \tag{24-4-2}$$

Variance

$$\begin{aligned}
 Var[X] &= E[(X - E[X])^2] \\
 &= E[X^2 - 2XE[X] + E^2[X]] \\
 &= E[X^2] - 2E[XE[X]] + E^2[X] \\
 &= E[X^2] - 2E^2[X] + E^2[X] \\
 &= E[X^2] - E^2[X]
 \end{aligned} \tag{24-4-3}$$

When X and Y are independent random variables,

$$Var[X + Y] = Var[X] + Var[Y]$$

In general, if n random variables

$$X_1, X_2, \dots, X_n$$

are pairwise independent, then

$$\text{Var}\left[\sum_{i=1}^n X_i\right] = \sum_{i=1}^n \text{Var}[X_i]$$

这个地方
不理解为什么用 pair-
wise, 而不用 mutually
independent

24.4.3 Indicator random variable

A random variable that has the value 1 or 0, according to whether a specified event occurs or not is called an indicator random variable for that event.

Handy facts: Suppose X is an indicator random variable for the event A . Let p denote $P(A)$.

Then

$$E(X) = p \quad (3.42)$$

$$\text{Var}(X) = p(1 - p) \quad (3.43)$$

These two facts are easily derived. In the first case we have, using our properties for expected value,

$$EX = 1 \times P(X = 1) + 0 \times P(X = 0) = P(X = 1) = P(A) = p \quad (3.44)$$

The derivation for $\text{Var}(X)$ is similar (use (3.29)).

The indicator random variable IA associated with event A has value 1 if event A occurs and has value 0 otherwise. In other words, IA maps all outcomes in the set A to 1 and all outcomes outside A to 0.

Random variables can be used to define events. In particular, any predicate involving random variables defines the event consisting of all outcomes for which the predicate is true.

e.g. for random variables R_1, R_2 , $R_1 = 1$ is an event, $R_2 = 2$ is an event, $R_1 = 1$ and $R_2 = 2$ is an event.